

Computer Systems

CS107

Cynthia Lee

Today's Topics

LAST TIME:

- Number representation
 - › Integer representation
 - › Signed numbers with two's complement

THIS TIME:

- Number representation
 - › The integer number line for signed and unsigned
 - › Overflow and underflow
 - › Comparison, extension and truncation in signed and unsigned
 - › Bitwise operations and bit sets

COMING UP:

- Today is last day of topics that will be included on next week's midterm
 - › Practice exams and topics list are up now

Signed integers with two's complement representation

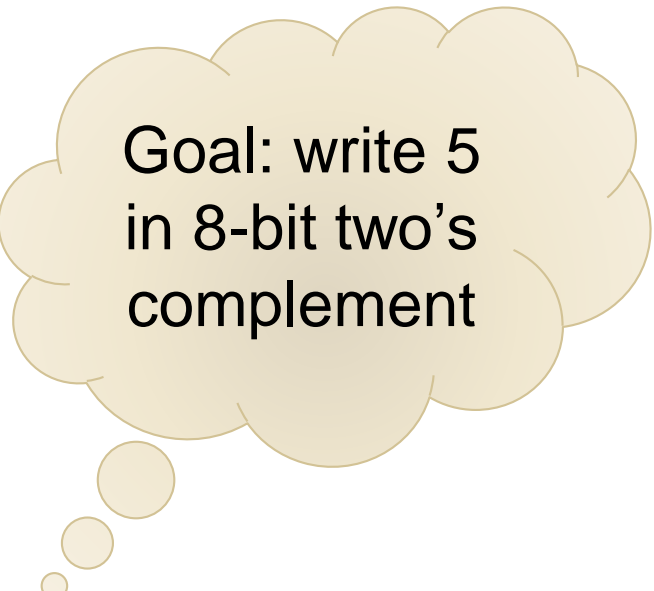
REVIEW FROM FRIDAY

Self-test review from Friday: Two's complement

What is the base-10 equivalent of the signed (two's complement) 4-bit binary number 1010?

- a) -10
- b) 10
- c) 11
- d) -11
- e) 5
- f) -5
- g) 6
- h) -6
- i) Other

Signed integers with two's complement

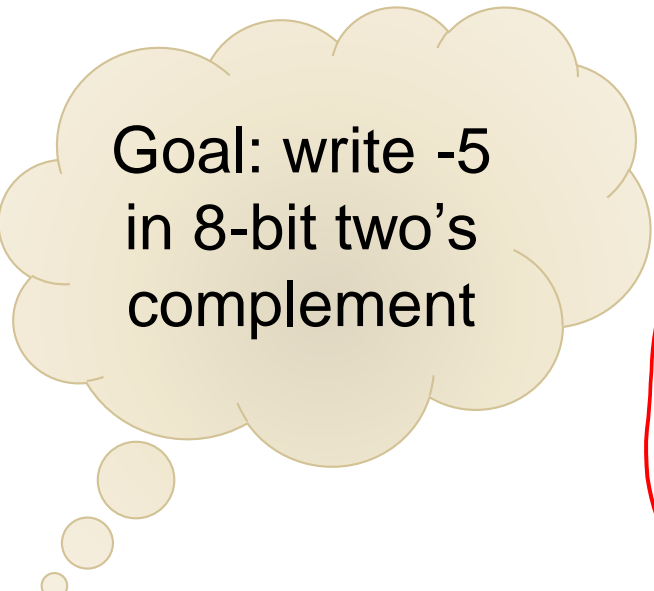


Goal: write 5
in 8-bit two's
complement

Steps to write a positive (or zero) number in two's complement:

1. Write the number in usual unsigned binary representation
2. Make sure that the number will “fit” in the number of bits you have
 - › For positive numbers, there needs to be at least one zero in the most significant (leftmost) bit
 - › 00000101 (no problem for 5 in 8 bits)
3. Done!
 - › Answer: 00000101

Signed integers with two's complement



Goal: write -5
in 8-bit two's
complement

Steps to write a negative number in two's complement:

1. Write the *absolute value* of the number in usual unsigned binary representation
2. Make sure that the number will “fit” in the number of bits you have
 - › Since we are writing the absolute value, a positive number, there needs to be at least one zero in the most significant (leftmost) bit*
 - › 00000101 (no problem for 5 in 8 bits)
3. “Flip” each bit ($0 \rightarrow 1$, $1 \rightarrow 0$)
 - › 00000101 \rightarrow 11111010
4. Add one
 - › 11111010 \rightarrow 11111011
5. Done!
 - › Answer: 11111011

* There is one negative number whose positive number won't “fit”—more on this Friday

What is the maximum value for a int32?



I can never remember that number. I need a memory rule.

606

integer



share improve this question



92

edited May 28 '14 at 14:09



Ben Hoffstein

49.5k 5 66 101

asked Sep 18 '08 at 17:18



Flinkman

5,181 4 18 48

107 Why would you need the exact number? I remember " $2^{31}-1$ " or " ± 2 billion" and that's good enough for everything I ever needed. – Joachim Sauer Mar 3 '09 at 11:21

27 unsigned: $2^{32}-1 = 4 \cdot 1024^3 - 1$; signed: $-2^{31} \dots +2^{31}-1$, because the sign-bit is the highest bit. Just learn $2^0=1$ to $2^{10}=1024$ and combine. $1024=1k$, $1024^2=1M$, $1024^3=1G$ – comonad Mar 28 '11 at 20:01

6 I generally remember that every 3 bits is about a decimal digit. This gets me to the right order of magnitude: 32 bits is 10 digits. – Barmar Oct 2 '13 at 15:11

30 Answers

active

oldest

votes



It's 2,147,483,647. Easiest way to memorize it is via a tattoo.

2397

share improve this answer



edited Oct 20 '14 at 16:30



Allbite

1,415 1 13 15

answered Sep 18 '08 at 17:20



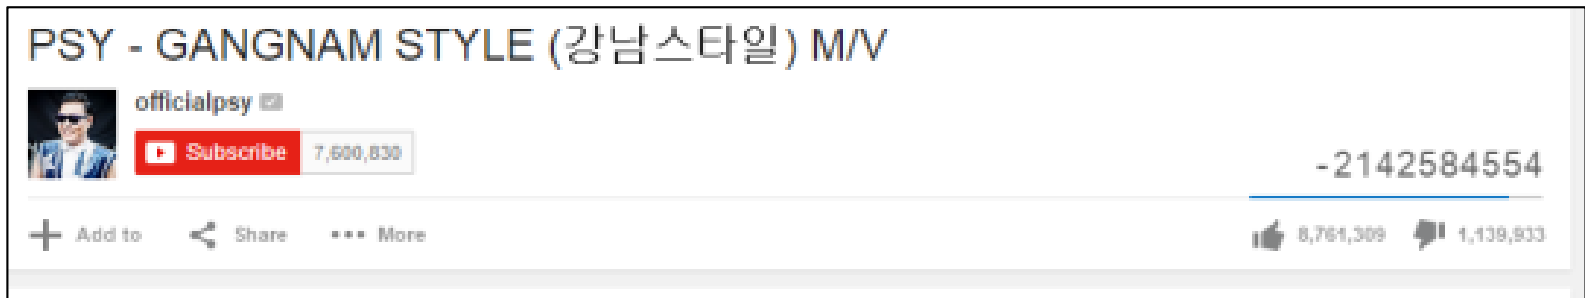
Ben Hoffstein

49.5k 5 66 101




Overflow in two's complement


- In two's complement, when you exceed the maximum value of int (2,147,483,647), you “wrap around” to negative numbers:



- Here is the link after Google upgraded to 64-bit integers:



Reasoning about signed and unsigned




Signed and unsigned numbers

At which points can overflow occur for signed and unsigned int?

(assume binary values shown are all 32 bits)

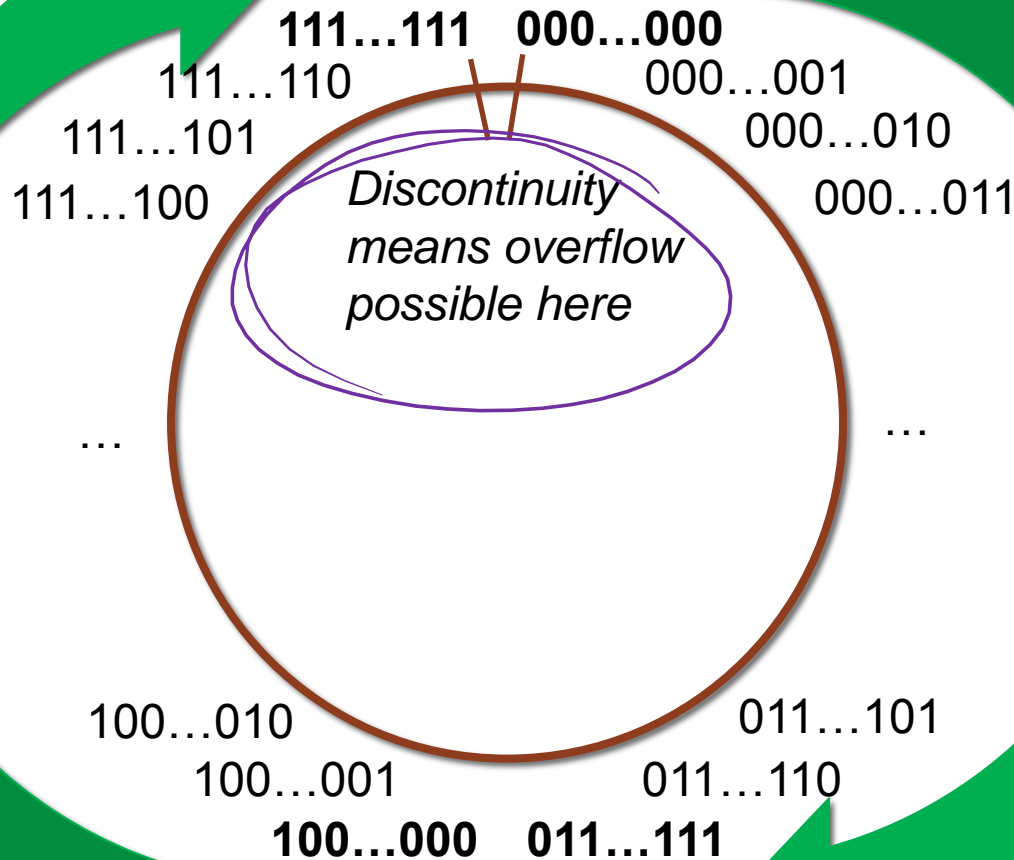
- A. Signed and unsigned can both overflow at points X and Y
- B. Signed can overflow at X, unsigned at Y
- C. Signed can overflow at Y, unsigned at X
- D. Signed can overflow at X and Y, unsigned only at X
- E. Other





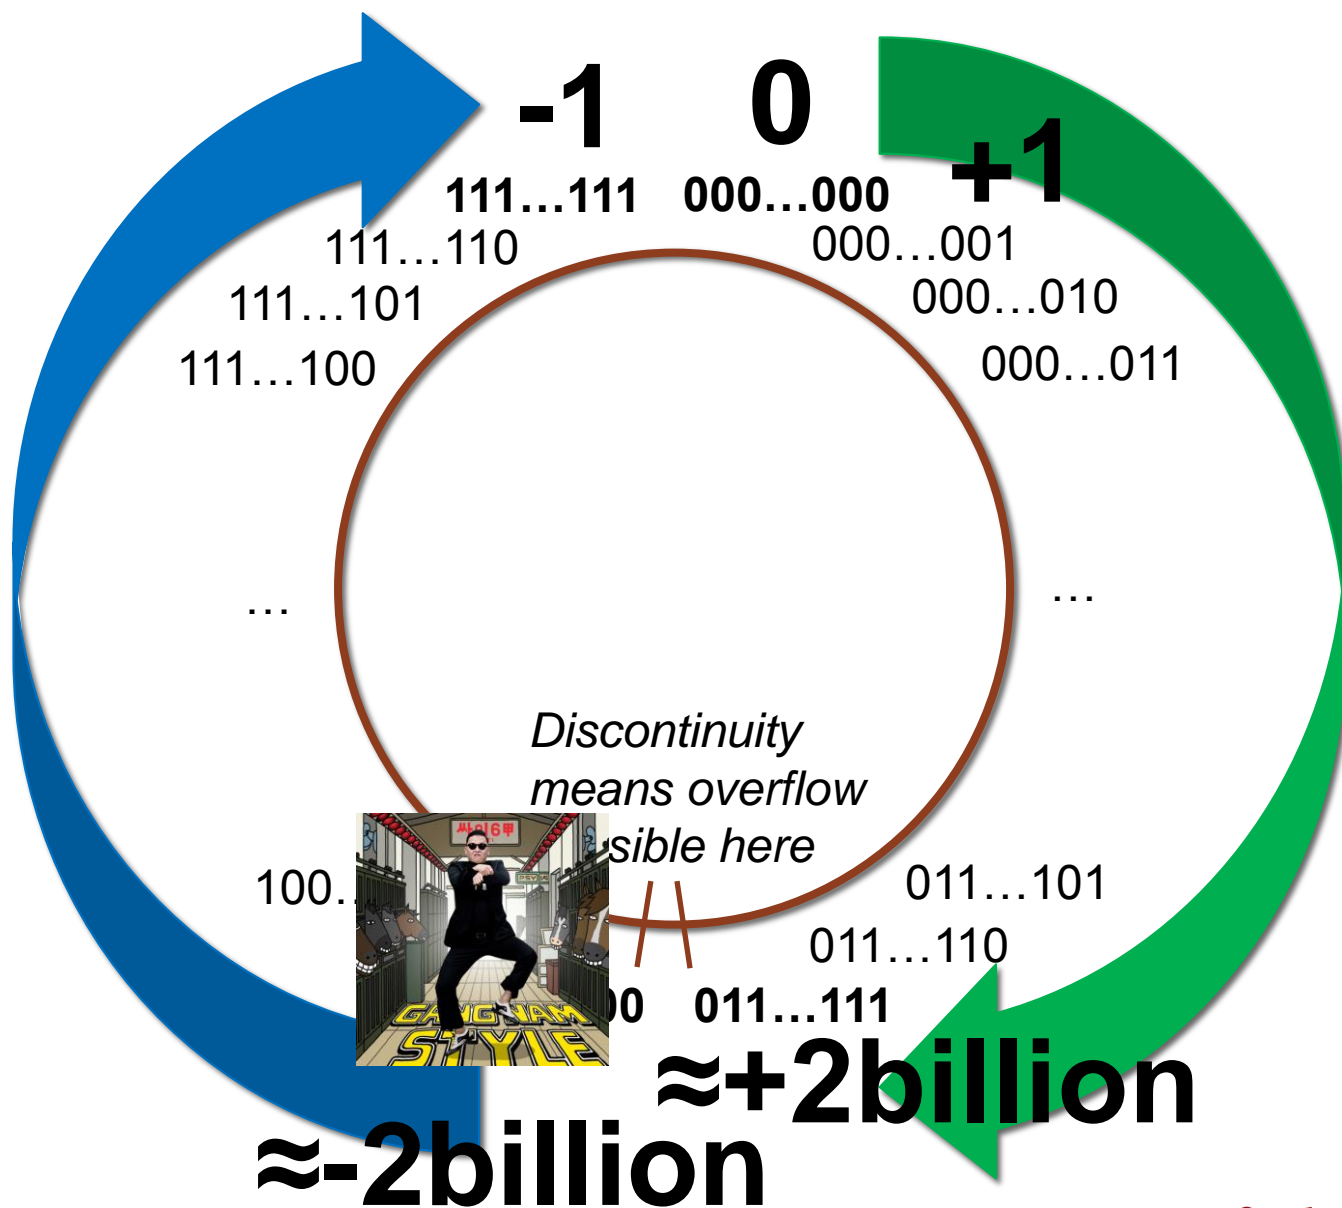
More increasing positive numbers

$\approx +4\text{billion } 0$



Increasing positive numbers

Negative numbers becoming less negative (i.e. increasing)



Increasing positive numbers

Comparison operators in signed and unsigned numbers

```
int      s1, s2, s3;
unsigned int u1, u2, u3;
```

Are the following statements true?


(assume that variables are set to values that place them in the spots shown)

> s3 > u3	Easy: true
> s1 > s3	Easy: false
> u1 > u3	Easy: true
> s1 > u3	Hmmm!??!

C just needs to choose one or the other scheme to dominate. It chooses...drumroll...

unsigned!

So this is **TRUE**.




HOME SELF-TEST:

Comparison operators in signed and unsigned numbers

```
int          s1, s2, s3, s4;  
unsigned int u1, u2, u3, u4;
```

Which many of the following statements are true? *(assume that variables are set to values that place them in the spots shown)*

- > s3 > u3
- > u2 > u4
- > s2 > s4
- > s1 > s2
- > u1 > u2
- > s1 > u3



Type truncation in the char/short/int/long family


```
int i1 = 0x8000007F; // = -2147483521
int i2 = 0x000000FF; // = 255
char s1 = i1; // = 0x7F = 127
char s2 = i2; // = 0xFF = -1
unsigned char u1 = i1; // = 0x7F = 127
unsigned char u2 = i2; // = 0xFF = 255
```

- Regardless of source or destination signed/unsigned type, truncation always just truncates
- This can cause the number to **change drastically in sign and value**

Type promotion in the char/short/int/long family

char	<code>sc = 0xFF;</code>	<code>// 0xFF = -1</code>
unsigned char	<code>uc = 0xFF;</code>	<code>// 0xFF = 255</code>
int	<code>s1 = sc;</code>	<code>// 0xFFFFFFFF = -1</code>
int	<code>s2 = uc;</code>	<code>// 0x000000FF = 255</code>
unsigned int	<code>u1 = sc;</code>	<code>// 0xFFFFFFFF = 4,294,967,295</code>
unsigned int	<code>u2 = uc;</code>	<code>// 0x000000FF = 255</code>

- Promotion always happens according to the *source* variable's type
 - › Signed: **“sign extension”** (copy MSB—0 or 1—to fill new space)
 - › Unsigned: **“zero fill”** (copy 0's to fill new space)
- *Note:* When doing <, >, <=, >= comparison between different size types, it will promote to the larger type
 - › “int < char” comparison will implicitly (1) assign char to int according to these promotion rules, *then* (2) do “int < int” comparison



Every base is base 10.

In closing

Bits As Individual Booleans

THIS IS A VERY DIFFERENT WAY OF THINKING ABOUT WHAT A PARTICULAR SET OF 8 BITS (ONE CHAR) “MEANS”

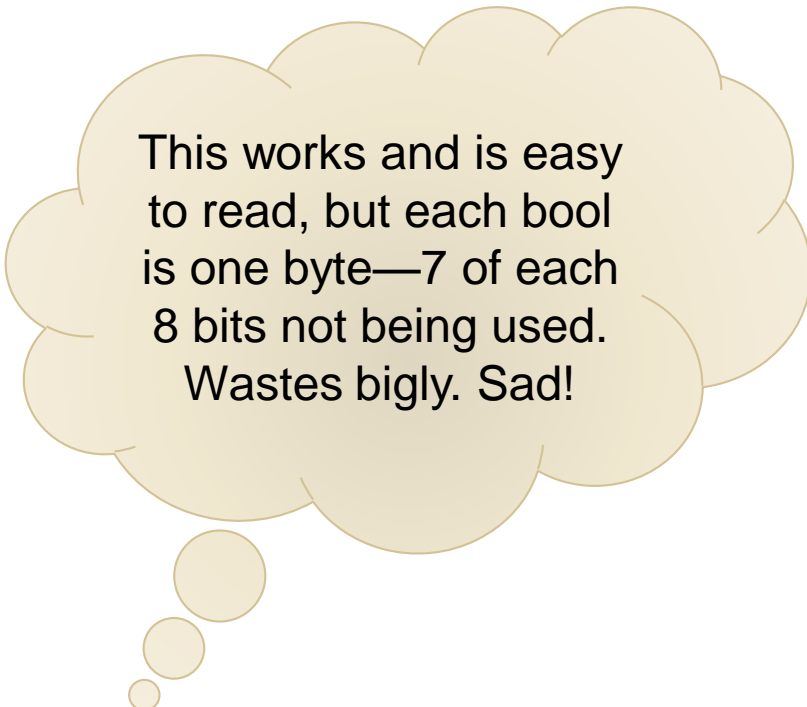
Bitwise operators and masking

- Let's say we want to represent font settings:
 - › **Bold**
 - › *Italic*
 - › Red color
 - › Superscript
 - › Underline
 - › ~~Strikethrough~~
- Observe that a particular piece of text can be any combination of these
 - › Example 1: ***Red***
 - › Example 2: *Red Underline*
 - › Example 3: ~~**^{Superscript} Underline Strikethrough**~~

Bitwise operators and masking

- Idea: Have a bool for each of these settings, store them in struct:

```
struct font_settings {  
    bool is_bold;  
    bool is_italic;  
    bool is_red;  
    bool is_super;  
    bool is_under;  
    bool is_strike;  
};
```



This works and is easy to read, but each bool is one byte—7 of each 8 bits not being used. Wastes bigly. Sad!

› Example 1: ***Bold Italic Red***

```
struct font_settings ex1; /* how to set up */  
ex1.is_bold = ex1.is_italic = ex1.is_red = true;  
ex1.is_super = ex1.is_under = ex1.is_strike = false;  
if (ex1.is_bold) { ...      /* how to use */
```

Bitwise operators and masking

- New idea: Have one 0/1 bit for each of these settings:
 - › **Bold** 1 = bold, 0 = not bold
 - › *Italic* 1 = italic, 0 = not italic
 - › **Red color** 1 = red, 0 = not red
 - › Superscript ...
 - › Underline
 - › ~~Strikethrough~~
- Store the collection of 6 bit settings together:
 - › Example 1: ***Red*** 111000
 - › Example 2: *Red Underline* 011010
 - › Example 3: ~~Bold Superscript Underline Strikethrough~~ 000111
- We can pack these into an unsigned char (uses lower 6 of the 8 bits)
 - › Example 1: ***Red*** 00111000

Bitwise operators and masking

- Use char and hexadecimal to store font settings:

Example 1: ***Red***

unsigned char ex1 = 0x38; // 0x38 = 00111000

Handwritten notes: 3 | 8

- ...But how do we use this?
- No way to “name” the bold bit by itself:**
if (ex1) { ... // tests if whole char != 0
if (ex1.is_bold) { ... // no nameable fields in char
- Can't access individual bits (system is byte-addressable)**
- Not hopeless: we need bitwise operators

Bitwise operators and bits as individual booleans

MOVING BEYOND THE “INT” INTERPRETATION OF BITS

Bitwise operators

- You've seen these categories of operators in C/C++:

- › Arithmetic operators: +, -, *, /
- › Comparison operators: ==, !=, <, >, <=, >=
- › Logical operators: &&, ||, !
- › (C++ only) Stream insertion operators: <<, >>

- Now meet a new category:

- › Bitwise operators: &, |, ^, ~, >>, <<

↑ ↑
shift operators

Bitwise operators

unsigned char a =		0	0	1	1	1	1	0	0
unsigned char b =		0	1	0	1	1	0	1	0
and, intersection	a & b	0	0	0	1	1	0	0	0
or, union	a b	0	1	1	1	1	1	1	0
xor, different?	a ^ b	0	1	1	0	0	1	1	0
not	~a	1	1	0	0	0	0	1	1
shift left ?	a << 2	1	1	1	1	0	0	0	0
shift right	a >> 3	0	0	0	0	0	1	1	1

Bitwise operators and masking

- Use char and hexadecimal to store font settings:

Example 1: ***Red***

```
unsigned char ex1 = 0x38;           // 0x38 = 00111000
```

- **How can we write a test for bold?**

```
bool is_bold(unsigned char settings)
{
    unsigned char mask = 1 << 5;    // 00100000
    return (mask & settings) != 0;
}
```

- **“Mask”** is what we call a number that we create solely for the purpose of extracting selected bits out of a bitwise representation
 - › Often crafted using 1 shifted by some amount
 - › Writing as a hexadecimal value also acceptable (0x20)
 - › More complex masks can be crafted in steps with | & etc to test for more than one condition at once

Bitwise operators and masking

- Reminder: here are our font settings, in bit order:
 - › **Bold**
 - › *Italic*
 - › Red color
 - › Superscript
 - › Underline
 - › ~~Strikethrough~~
- How can we write code to turn off italics (*without* changing any other settings)?

```
unsigned char italics_off(unsigned char settings)
{
    return _____;
}
```

A. ~settings

B. settings & (1 << 4)

C. settings ^ (1 << 4)

D. settings | (~(1 << 4))

E. settings & (~(1 << 4))

F. Something else

(to be) || !(to be), that is the question

- ! and ~ are both “not” operators—are they the same?
- **In other words, is this guaranteed to always print?**

```
int i;  
scanf("%d", &i);  
if ((!i) == (~i)) printf("same this time\n");
```

A. Yes, always prints

☒ B. Sometimes prints, but not always

C. No, never prints

D. You lost me at the code version of Shakespeare

i = -1 is special case

$!(-1) \Rightarrow 0$ $!(3) \Rightarrow 0$
 $\sim(-1) \Rightarrow 0000\dots000$ $\sim(3) \Rightarrow 111\dots1100$