

Computer Systems

CS107

Cynthia Lee

Today's Topics

LECTURE:

- › Floating point!
- › (if we have time)
 - Preview of Friday's assembly language lecture

Real Numbers and Approximation

MATH TIME!

Some preliminary observations on approximation

- We know that some non-integer numbers can be represented precisely by a finite-length decimal value
 - › $1/5_{10} = 0.2_{10}$
- ...and some can't (in base 10)
 - › π
 - › $1/3_{10} = .3333333\overline{3}_{10}$
- But this isn't consistent across different bases
 - › $1/3_{10} = 0.1_3$
 - › $1/5_{10} = 0.0120120\overline{12}_3$
- Uncountably many real numbers we could represent!
- And of course even with `int(countable)` we needed to make choices about which integers make the cut and which don't, when *we have only a limited storage space*



Real number types: float, double

- If you're deciding how to divide up a large number of possible bit patterns into numbers they can represent, there are any number of ways you could do this
 - › What are your priorities?
 - › Do you want to leave space for very big ones? Or very small ones?
 - › Do you prefer positive or negative?



Thought experiment: how we could represent real numbers

1. Take our integer representation
 2. Decide how much granularity we want to represent
 3. Move the decimal point (binary point?) over a corresponding number of places
- For example: let's say we want to be able to measure in units of $1/64^{\text{th}}$
 - › 32-bit binary number example (10 and $1/64^{\text{th}}$):
 - 0000 0000 0000 0000 0000 0010 10.00 0001



Real number types: float, double

- **float** is 4 bytes = 32 bits
 - $2^{32} = 4,294,967,296$ possible values
 - (note: obviously same number of possible values as 32-bit int)
- **double** is 8 bytes = 64 bits
 - $2^{64} = 18,446,744,073,709,551,616$ possible
- This is the same number of different possible representable values as int and long, of course
 - › But floating point range goes much higher than max int:
 - › Max 32-bit int: $2^{31} - 1 = 2,147,483,647 \approx 2.1 \times 10^9$
 - › Max 32-bit float: $(2 - 2^{-23}) \times 2^{127} \approx 3.4 \times 10^{38}$



The “mini-float”

NOT A REAL TYPE IN C, BUT WILL SHOW US THE PRINCIPLES THAT
WILL SCALE UP TO THE ACTUAL BIT SIZES FLOAT AND DOUBLE

IEEE format, squished down to example-size “mini-float”




- **Sign** bit is 1 for negative, 0 for positive
 - › Operates completely independently, *unlike* 2's complement int
- The rest of the number is in something like scientific notation:
 - › $(+/-) M * 2^{\text{Exp}}$
 - **Mantissa** provides the significant digits of the number
 - *Note:* also called the coefficient or the significand
 - **Exponent** is used to scale the number up or down to very large or very small values
 - › Both exponent and mantissa are binary numbers, 4 bits and 3 bits, respectively
 - *but not in normal 2's complement form!*

IEEE format, squished down to example-size “mini-float”



- Compare to the bit distribution for 32-bit float :



Creative Commons license attribution to user Fresheneesz at
https://en.wikipedia.org/wiki/Single-precision_floating-point_format#/media/File:Float_example.svg

“Mini-float”: the mantissa

- A 3-bit binary number representing a number after an implicit 1.
 - › What??
- Yes, it looks like this: $1.\textit{[mantissa digits]}$
- Examples:
 - › Mantissa = $100_2 \rightarrow 1.100_2$
 - $= 1 + \frac{1}{2} = 1.5_{10}$
 - › Mantissa = $110_2 \rightarrow 1.110_2$
 - $= 1 + \frac{1}{2} + \frac{1}{4} = 1.75_{10}$

Sign	Exponent				Mantissa		

“Mini-float”: the exponent

- A 4-bit binary number representing the exponent:

› 0000	› reserved
› 0001	› -6
› 0010	› -5
› 0011	› -4
› 0100	› -3
› 0101	› -2
› 0110	› -1
› 0111	› 0
› 1000	› 1
› 1001	› 2
› 1010	› 3
› 1011	› 4
› 1100	› 5
› 1101	› 6
› 1110	› 7
› 1111	› reserved

Sign	Exponent				Mantissa		


Lets do an example!

Sign	Exponent				Mantissa		
0	0	0	0	1	0	1	0

- This number is:
 - A. Greater than 0
 - B. Less than 0
 - C. Help!
- Extra credit: > 1 or < -1 or in between?


Lets do an example!

+/-	Exponent				Mantissa			Value
0	0	0	0	1	0	1	0	$10/8 * 2^{-6} = 10/512$
0	0	0	0	1	0	1	1	$11/8 * 2^{-6} = 11/512$
0	0	0	0	1	1	0	0	$12/8 * 2^{-6} = 12/512$
0	0	0	0	1	1	0	1	$13/8 * 2^{-6} = 13/512$
0	0	0	0	1	1	1	0	$14/8 * 2^{-6} = 14/512$
0	0	0	0	1	1	1	1	$15/8 * 2^{-6} = 15/512$
0	0	0	1	0	0	0	0	$1 * 2^{-5} = 16/512$
0	0	0	1	0	0	0	1	$9/8 * 2^{-6} = 18/512$



Lets do an example!

+/-	Exponent				Mantissa			Value
0	0	0	1	0	0	0	0	$1 * 2^{-5} = 16/512$
0	0	0	1	0	0	0	1	$9/8 * 2^{-6} = 18/512$
0	1	1	1	0	0	1	0	$10/8 * 2^7 = 160$
0	1	1	1	0	0	1	1	$11/8 * 2^7 = 176$
0	1	1	1	0	1	0	0	$12/8 * 2^7 = 192$
0	1	1	1	0	1	0	1	$13/8 * 2^7 = 208$
0	1	1	1	0	1	1	0	$14/8 * 2^7 = 224$
0	1	1	1	0	1	1	1	$15/8 * 2^7 = 240$



About those reserved exponents...

DENORMALIZED AND SPECIAL CASES

Reserved exponent values: 0000 and 1111

0000 exponent:

- › If the mantissa is all zeros: ± 0

Sign	Exponent				Mantissa		
any	0	0	0	0	0	0	0

- › If the mantissa is nonzero: denormalized floats

Sign	Exponent				Mantissa		
any	0	0	0	0	any nonzero		

1111 exponent:

- › If the mantissa is all zeros: $\pm \text{INF}$ (infinity)

Sign	Exponent				Mantissa		
any	1	1	1	1	0	0	0

- › If the mantissa is nonzero: NaN (not a number)

Sign	Exponent				Mantissa		
any	1	1	1	1	any nonzero		

MORE about: the mantissa

- A 3-bit binary number representing a number after an implicit 1.
 - › What??
- Yes, it looks like this: $1.[mantissa\ digits]$
- Examples:
 - › Mantissa = $100_2 \rightarrow 1.100_2$
 - $= 1 + \frac{1}{2} = 1.5_{10}$
 - › Mantissa = $110_2 \rightarrow 1.110_2$
 - $= 1 + \frac{1}{2} + \frac{1}{4} = 1.75_{10}$
- **If the exponent is all zeros, then there is no implicit leading 1**
 - › The exponent is implicitly **smallest possible exponent** (2^{-6} for mini-float)
 - › This allows us to eke out a few numbers that are even (closer to zero) than would otherwise be possible
 - › Called “denormalized” or “denorm” floats

Sign	Exponent				Mantissa		

Comparing float and int

ADVANTAGES AND DISADVANTAGES

Comparing float and int

- 32-bit integer (type **int**):
 - › -2,147,483,648 to 2147483647
 - › Every integer in that range can be represented
- 64-bit integer (type **long**):
 - › -9,223,372,036,854,775,808 to 9,223,372,036,854,775,807
 - (Even Psy will have a hard time overflowing this!)
- 32-bit floating point (type **float**):
 - › $\sim 1.7 \times 10^{-38}$ to $\sim 3.4 \times 10^{38}$
 - › Not all numbers in the range can be represented (obviously—uncountable)
 - Not even all integers in the range can be represented!
 - Gaps can get quite large!! (larger the exponent, larger the gap between successive mantissa values)
- 64-bit floating point (type **double**):
 - › $\sim 9 \times 10^{-307}$ to $\sim 2 \times 10^{308}$

Doing arithmetic in float

FUN TIMES

Adding two mini-floats

- Your bank account balance on your 12th birthday was \$128.00
- Your bank stores this amount as a mini-float:

Sign	Exponent				Mantissa		

- Each week of your childhood, starting at age 12, you deposited your weekly allowance of \$8:

Sign	Exponent				Mantissa		

- **What was your account balance on the day you turned 18?**
 - A. About \$10000
 - B. About \$5000
 - C. About \$2500
 - D. About \$100
 - E. Other

Doing arithmetic in float

FUN TIMES

Adding two mini-floats

- Your bank account balance on your 12th birthday was \$128.00
- **Your bank stores this amount as a mini-float:**

Sign	Exponent				Mantissa		

- Each week, starting at age 12, you deposited your weekly allowance of \$14:

Sign	Exponent				Mantissa		

- **What was your account balance on the day you turned 18?**
 - A. About \$10000
 - B. About \$5000
 - C. About \$2500
 - D. About \$100
 - E. Other

Instruction Set Architectures

SOME CONTEXT AND TERMINOLOGY

Instruction Set Architecture

- **The ISA defines:**

- › Operations that the processor can execute
- › Data transfer operations + how to access data
- › Control mechanisms like branch, jump (think loops and if-else)
- › Contract between programmer/compiler and hardware

- **Layer of abstraction:**

- › Above:
 - Programmer/compiler can write code for the ISA
 - New programming languages can be built on top of the ISA as long as the compiler will do the translation
- › Below:
 - New hardware can implement the ISA
 - Can have even potentially radical changes in hardware implementation
 - Have to “do” the same thing from programmer point of view

- **ISAs have incredible inertia!**

- › Legacy support is a huge issue for x86-64

Two major categories of Instruction Set Architectures

- **CISC:**

- › **Complex** instruction set computers
 - e.g., x86 (**CS107 studies this**)
- › Have special instructions for each thing you might want to do
- › Can write code with fewer instructions, because each instruction is very expressive



- **RISC:**

- › **Reduced** instruction set computers
 - e.g., MIPS
- › Have only a very tiny number of instructions, optimize the heck out of them in the hardware
- › Code may need to be longer because you have to go roundabout ways of achieving what you wanted

