

Computer Systems

CS107


Cynthia Lee

Today's Topics

LECTURE:

- › Pop quiz on floating point!
 - Just kidding.
- › **New:** Assembly code

Two friendly reminders:





REMINDER:
Everything is bits!

Everything is bits!

- We've seen many data types so far:
 - › Integers:
 - char/short/int/long (encoding as unsigned or two's complement signed)
 - › Letters/punctuation/etc:
 - char (ASCII encoding)
 - › Real numbers:
 - float/double (IEEE floating point encoding)
 - › Memory addresses:
 - pointer types (unsigned long encoding)
 - › Now a new one.....**the code itself!**
 - Instructions (AMD64 encoding)



What happens when we compile our code?

ANATOMY OF AN EXECUTABLE FILE

What happens when we compile our code?

```
int sum_array(int arr[], int nelems) {  
    int sum = 0;  
    for (int i = 0; i < nelems; i++) {  
        sum += arr[i];  
    }  
    return sum;  
}
```

```
> make  
> ls  
Makefile  sum  sum.c  
> objdump -d sum
```

000000000040055d <sum_array>:

40055d:	ba 00 00 00 00	mov	\$0x0,%edx
400562:	b8 00 00 00 00	mov	\$0x0,%eax
400567:	eb 09	jmp	400572
400569:	48 63 ca	movslq	%edx,%rcx
40056c:	03 04 8f	add	(%rdi,%rcx,4),%eax
40056f:	83 c2 01	add	\$0x1,%edx
400572:	39 f2	cmp	%esi,%edx
400574:	7c f3	j1	400569
400576:	f3 c3	repz retq	

000000000040055d <sum_array>:

40055d: ba 00 00 00 00

mov \$0x0,%edx

400562: b8 00 00 00 00

mov \$0x0,%eax

400567

400572

400569

lq %edx,%rcx

40056c

(%rdi,%rcx,4),%eax

40056f

\$0x1,%edx

400572

%esi,%edx

400574: 7c f3

j1 400569

400576: f3 c3

repz retq

Name of the function (same as in the C code) and the memory address where the code for this function starts

000000000040055d <sum_array>:

40055d:	ba 00 00 00 00	mov	\$0x0,%edx
400562:	b8 00 00 00 00	mov	\$0x0,%eax
400567:	eb 09	jmp	400572
400569:	vsq	%edx,%rcx
40056c:		d	(%rdi,%rcx,4),%eax
40056f:		d	\$0x1,%edx
400572:		p	%esi,%edx
400574:			400569
400576:		pz retq	

Memory address
where each of line of
instruction is found—
sequential instructions
are found sequentially
in memory

000000000040055d <sum_array>:

```
40055d:    ba 00 00 00 00
400562:    b8 00 00 00 00
400567:    eb 09
400569:    48 63 ca
40056c:    03 04 0f
40056f:
400572:
400574:
400576:
```

Assembly code:
“human-readable”
version of each
instruction

```
mov    $0x0,%edx
mov    $0x0,%eax
jmp    400572
movslq %edx,%rcx
add    (%rdi,%rcx,4),%eax
add    $0x1,%edx
cmp    %esi,%edx
jl     400569
repz   retq
```

000000000040055d <sum_array>:

```
40055d:    ba 00 00 00 00
400562:    b8 00 00 00 00
400567:    eb 09
400569:    48 63 ca
40056c:    03 04 8f
40056f:    83 c2 01
400572:    39 f2
400574:    7c f3
400576:    f3 c3
```

```
mov    $0x0,%edx
mov    $0x0,%eax
jmp     400572
```

Machine code:
raw hexadecimal
version of each
instruction,
representing the
binary as it would
be read by the
computer


4),%eax

Anatomy of an individual instruction

Anatomy of an individual instruction

40056c:	03 04 8f	add	(%rdi,%rcx,4),%eax
40056f:	83 c2 01	add	\$0x1,%edx
400572:	39 f2	cmp	%esi,%edx
400574:	7c f3	j1	400569

Operation name
(sometimes called
“opcode”)




Operands
(like arguments)



Anatomy of an individual instruction

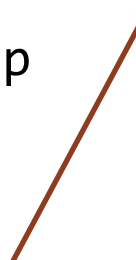
40056c:	03 04 8f	add	(%rdi,%rcx,4),%eax
40056f:	83 c2 01	add	\$0x1,%edx
400572:	39 f2	cmp	%esi,%edx
400574:	7c f3	j1	400569



%[name] names a register—these are a small collection of memory slots right on the CPU that can hold variables' values

Anatomy of an individual instruction

40056c:	03 04 8f	add	(%rdi,%rcx,4),%eax
40056f:	83 c2 01	add	\$0x1 ,%edx
400572:	39 f2	cmp	%esi,%edx
400574:	7c f3	j1	400569



\$[number] means a constant
value (this is the number 1)

add and sub instruction breakdowns

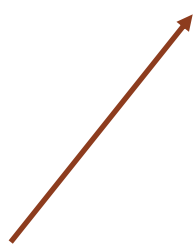
Op	Source1	Source2/Dest	Dest Comments
add	op1	op2	op2 += op1

Op	Source1	Source2/Dest	Dest Comments
sub	op1	op2	op2 += op1

- Note that you have no choice but to overwrite op2 with the sum (or difference) of op2 and op1
 - › *No separate third destination operand*
- Op1 and op2 can be registers (e.g., %rax) or constants (e.g., \$1)
 - › *Plus some more complex options we will learn about later*

000000000040055d <sum_array>:

40055d:	ba 00 00 00 00	mov	\$0x0,%edx
400562:	b8 00 00 00 00	mov	\$0x0,%eax
400567:	eb 09	jmp	400572
400569:	48 63 ca	movslq	%edx,%rcx
40056c:	03 04 8f	add	(%rdi,%rcx,4),%eax
40056f:	83 c2 01	add	\$0x1,%edx
400572:	39 f2	cmp	%esi,%edx
400574:	7c f3	j1	400569
400576:	f3 c3	repz	retq



Guessing game: which part of the `sum_array` C code do you think corresponds to the marked assembly instruction?

```
int sum_array(int arr[], int nelems) {  
    int sum = 0;                                // (A)  
    for (int i = 0; i < nelems; i++) {           // (B)  
        sum += arr[i];                           // (C)  
    }                                             // or (D) other?  
    return sum;  
}
```

```
<add_stuff>:  
mov    $0x0,%eax
```

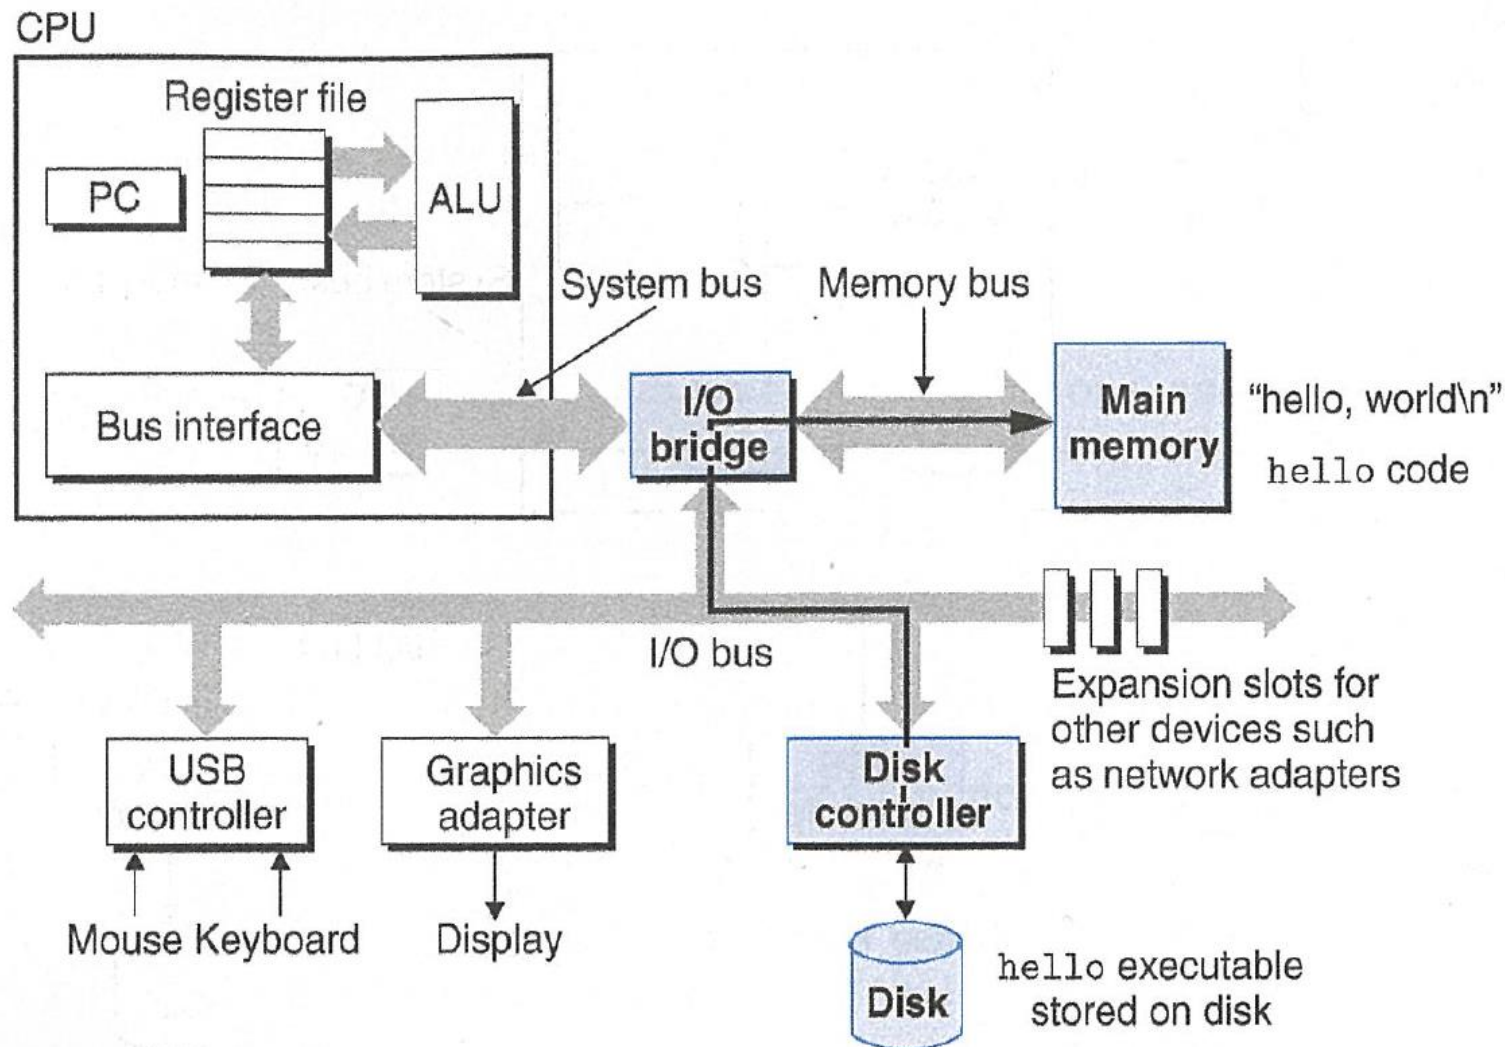
Your turn: can we write assembly code to match this C code?

```
int add_stuff(int edi, int esi, int edx) {  
    int sum = 0;  
    edi = edi + 3;  
    esi = esi + edx;  
    sum = edi + esi;  
    return sum;  
}
```

Registers and memory

ANATOMY OF THE COMPUTER

An architecture view of computer hardware



The mov instruction

OUR FIRST INSTRUCTION

Dude, where's my data?

- **A main job of assembly language is to manage data:**
 - › Data can be on the CPU (in registers) or in memory (at an address)
 - Turns out this distinction REALLY MATTERS for performance
 - https://people.eecs.berkeley.edu/~rcs/research/interactive_latency.html
 - › Instructions often want to move data:
 - Move from one place in memory to another
 - Move from one register to another
 - Move from memory to register
 - Move from register to memory
 - › Instructions often want to operate on data:
 - Add contents of register X to contents of register Y
- **Hence “mov” (move) instruction is paramount!**

mov

- **mov src, dst**
 - › Optional suffix (b,w,l,q): movb, movw, movl, movq
 - › One confusing thing about “move” it makes it sound like it leaves the src “empty”—no!
 - Does a **copy**, like the assignment operator you are familiar with
 - › src, dst options:
 - Immediate (AKA constant value)
 - Register
 - Memory

Basic addressing modes (Think: assembly version of VARIABLES)

- Notice that one major difference between high-level code and assembly instructions is the absence of programmer-chosen, descriptive variable names:

```
int total_goodness = nReeses + nButterfinger;
```

```
addl 8(%rbp),%eax
```

- We don't get to choose variable names, we have to talk directly about places in hardware
- **“Addressing modes” are allowable ways of naming these places**

Basic addressing modes (Think: assembly version of VARIABLES)

Op	Source	Dest	Dest Comments
movl	\$0,	%eax	Name of a register
movl	\$0,	0x8f2713e0	Actual address literal (note address literals are different from other literals—don't need \$ in front)
movl	\$0,	(%rax)	Look in the register named, find an address there, and use it

Reminder: need to put \$ in front of immediate values (constant literals)

Basic addressing modes (Think: assembly version of VARIABLES)

Op	Source	Dest	Dest Comments
movl	\$0,	%eax	Name of a register
movl	\$0,	0x8f2713e0	Actual address literal (note address literals are different from other literals—don't need \$ in front)
movl	\$0,	(%rax)	Look in the register named, find an address there, and use it
movl	\$0,	-24(%rbp)	Add -24 to an address in the named register, and use that address

Displacement must be a constant; to have a variable base and variable displacement, use two steps: add then mov

Displacement can be positive or negative

Basic addressing modes (Think: assembly version of VARIABLES)

Op	Source	Dest	Dest Comments
movl	\$0,	%eax	Name of a register
movl	\$0,	0x8f2713e0	Actual address literal (note address literals are different from other literals—don't need \$ in front)
movl	\$0,	(%rax)	Look in the register named, find an address there, and use it
movl	\$0,	-24(%rbp)	Add -24 to an address in the named register, and use that address
movl	\$0	8(%rbp, %eax, 2)	Address to use = (8 + address in rbp) + (2 * index in eax)

Any constant allowed

Only 1, 2, 4, 8 allowed

Matching exercise: Addressing modes use cases

- Match up which use cases make the most sense for which addressing modes (some guesswork expected)

Op	Src	Dest	Use case?
movl	\$0	8(%rbp, %eax, 2)	
movl	\$0,	%eax	
movl	\$0,	0x8f2713e0	
movl	\$0,	4(%rax)	

Use cases	
(a)	Prepare to use 0 in a calculation
(b)	Zero out a field of a struct
(c)	Zero out a given array bucket
(d)	Zero out a global variable

Instruction Set Architectures

SOME CONTEXT AND TERMINOLOGY

Instruction Set Architecture

- **The ISA defines:**

- › Operations that the processor can execute
- › Data transfer operations + how to access data
- › Control mechanisms like branch, jump (think loops and if-else)
- › Contract between programmer/compiler and hardware

- **Layer of abstraction:**

- › Above:
 - Programmer/compiler can write code for the ISA
 - New programming languages can be built on top of the ISA as long as the compiler will do the translation
- › Below:
 - New hardware can implement the ISA
 - Can have even potentially radical changes in hardware implementation
 - Have to “do” the same thing from programmer point of view

- **ISAs have incredible inertia!**

- › Legacy support is a huge issue for x86-64

Two major categories of Instruction Set Architectures

▪ CISC:

- › **Complex** instruction set computers
 - e.g., x86 (**CS107 studies this**)
- › Have special instructions for each thing you might want to do
- › Can write code with fewer instructions, because each instruction is very expressive



▪ RISC:

- › **Reduced** instruction set computers
 - e.g., MIPS
- › Have only a very tiny number of instructions, optimize the heck out of them in the hardware
- › Code may need to be longer because you have to go roundabout ways of achieving what you wanted

