

Computer Systems

CS107

Cynthia Lee

Today's Topics

LECTURE:

- › **More** assembly code!
 - Arithmetic and logic operations
 - If-else control
 - Loops (to be continued Friday)

Basic addressing modes (Think: assembly version of VARIABLES)

Op	Source	Dest	Dest Comments
movl	\$0,	%eax	Name of a register
movl	\$0,	0x8f2713e0	Actual address literal (note address literals are different from other literals—don't need \$ in front)
movl	\$0,	(%rax)	Look in the register named, find an address there, and use it
movl	\$0,	<u>-24</u> (%rbp)	Add -24 to an address in the named register, and use that address
movl	\$0	8(%rbp, %eax, 2)	Address to use = (8 + address in rbp) + (2 * index in eax)

Reference material: arithmetic and bitwise operations

WE ALREADY SAW **ADD** AND **SUB**, HERE ARE MORE!

Arithmetic and bitwise operations

<code>add src, dst</code>	<code># dst += src</code>
<code>sub src, dst</code>	<code># dst -= src</code>
<code>imul src, dst</code>	<code># dst *= src</code>
<code>inc dst</code>	<code># dst += 1</code>
<code>dec dst</code>	<code># dst -= 1</code>
<code>neg dst</code>	<code># dst = -dst</code>
<code>and src, dst</code>	<code># dst &= src</code>
<code>or src, dst</code>	<code># dst = src</code>
<code>xor src, dst</code>	<code># dst ^= src</code>
<code>not dst</code>	<code># dst = ~dst</code>
<code>shl count, dst</code>	<code># dst <<= count (left shift)</code>
<code>sar count, dst</code>	<code># dst >>= count (arithmetic right shift)</code>
<code>shr count, dst</code>	<code># dst >>= count (logical right shift)</code>

Reference material: How to view assembly on myth

HANDY TIPS

Objdump: makes readable assembly from executable

- Gives you output like below, for all functions in executable
- Doesn't execute or debug code, just provides this output for you to read

```
myth5> objdump -d sum // replace "sum" with your executable
```

```
000000000040055d <sum_array>:
```

40055d:	ba 00 00 00 00	mov	\$0x0,%edx
400562:	b8 00 00 00 00	mov	\$0x0,%eax
400567:	eb 09	jmp	400572
400569:	48 63 ca	movslq	%edx,%rcx
40056c:	03 04 8f	add	
	(%rdi,%rcx,4),%eax		
40056f:	83 c2 01	add	\$0x1,%edx
400572:	39 f2	cmp	%esi,%edx
400574:	7c f3	j1	400569
400576:	f3 c3	repz retq	

Gdb: debug in C or assembly

- In addition to all its other wonderful features, gdb lets you see assembly version of code you are debugging
 - › You can step through code either line-by-line in C code, or instruction-by-instruction in assembly code
 - › ...Or see both at the same time!

myth5> gdb sum // replace "sum" with your executable

Reading symbols from sum...done.

(gdb) break main

(gdb) run

(gdb) layout split B+>

```
39
40      int main(int argc, char *argv[])
41
0x400578 <main>          sub     $0x28,%rsp
0x40057c <main+4>         movl    $0x58,(%rsp)
0x400583 <main+11>        ("Sum movl    $0x5e,0x4(%rsp)res, nscores));
0x40058b <main+19>        movl    $0x46,0x8(%rsp)
0x400593 <main+27>        movl    $0x5c,0xc(%rsp)
0x40059b <main+35>        movl    $0x53,0x10(%rsp)bly Available ]
0x4005a3 <main+43>        movl    $0x5c,0x14(%rsp)
0x4005ab <main+51>        movl    $0x1,0x18(%rsp)
0x4005b3 <main+59>        movl    $0x5c,0x1c(%rsp)
0x4005bb <main+67>        mov     $0x8,%esi
0x4005c0 <main+72>        mov     %rsp,%rdi
B+>
child process 22274 In: main
```


Examples:
Let's read code!

PREPARATION FOR YOUR ASSIGN5

Preparing for assign5: reading assembly code

0000000004005ac <sum_example1>:

4005bd: 8b 45 e8

mov %esi,%eax

4005c3: 01 d0

add %edi,%eax

4005cc: c3

retq

// (A)

```
void sum_example1() {
    int x, y;
    int sum = x + y;
}
```

// (B)

```
void sum_example1(int x, int y) {
    int sum = x + y;
}
```

// (C)

```
int sum_example1() {
    int x, y;
    return x + y;
}
```

// (D)

```
int sum_example1(int x, int y) {
    return x + y;
}
```

Reference material: Register special uses

ALL REGISTERS ARE JUST BOXES IN HARDWARE, BUT SOME HAVE SPECIAL USES (SOME PURELY BY “TRADITION,” OTHERS DICTATED BY ASSEMBLY LANGUAGE WORKINGS)

Register uses (includes a few of the most common—for more complete list see reference on course website)

Register	Conventional use	Low 32-bits	Low 16-bits	Low 8-bits
%rax	Return value	%eax	%ax	%al
%rdi	1st argument	%edi	%di	%dil
%rsi	2nd argument	%esi	%si	%sil
%rdx	3rd argument	%edx	%dx	%dl
%r10	Scratch/temporary	%r10d	%r10w	%r10b
%r11	Scratch/temporary	%r11d	%r11w	%r11b
%rip	Instruction pointer			
%rflags	Status/condition code bits			

Preparing for assign5: reading assembly code

000000000400578 <sum_example2>:

400578:	8b 47 0c	mov	0xc(%rdi),%eax
40057b:	03 07	add	(%rdi),%eax
40057d:	2b 47 18	sub	0x18(%rdi),%eax
400580:	c3	retq	

```
int sum_example2(int arr[]) {
    int sum = 0;
    sum += arr[0];
    sum += arr[3];
    sum -= arr[6];

    return sum;
}
```

Which register or memory address represents the C variable sum?

- (a) 0xc(%rdi)
- (b) %rdi
- (c) (%rdi)
- (d) 0x18(%rdi)
- (e) %eax

Preparing for assign5: reading assembly code

000000000400578 <sum_example2>:

400578:	8b 47 0c	<i>sum = arr[3]</i>	mov	0xc(%rdi),%eax
40057b:	03 07	<i>sum += arr[0]</i>	add	(%rdi),%eax
40057d:	2b 47 18		sub	0x18(%rdi),%eax
400580:	c3		retq	

```
int sum_example2(int arr[]) {
    int sum = 0;
    sum += arr[0];
    sum += arr[3];
    sum -= arr[6];
    return sum;
}
```

Which register or memory address represents the C value 6 (as in arr[6])?

- (a) 0xc
- (b) %rdi
- (c) (%rdi)
- (d) 0x18
- (e) %eax

What does it mean for a program to execute?

HOW DO WE MOVE FROM ONE INSTRUCTION TO THE NEXT?

HOW DO COMPUTERS “DO STUFF”?

Data storage vs. “doing stuff” on a computer

- Data sits in memory (and, we now know, registers)
- We understand that there are instructions that control movement of the data and operations on it
- ...
- **But who controls the instructions? How do we know what to do now? ...or do next?**

4004fd	fa
4004fc	eb
4004fb	01
4004fa	fc
4004f9	45
4004f8	83
4004f7	00
4004f6	00
4004f5	00
4004f4	00
4004f3	fc
4004f2	45
4004f1	c7
4004f0	e5
4004ef	89
4004ee	48
4004ed	55



Progress through a function

00000000004004ed <loop>:

→ 4004ed:	55	push	%rbp
4004ee:	48 89 e5	mov	%rsp,%rbp
4004f1:	c7 45 fc 00 00 00 00	movl	\$0x0,-0x4(%rbp)
4004f8:	83 45 fc 01	addl	\$0x1,-0x4(%rbp)
4004fc:	eb fa	jmp	4004f8 <loop+0xb>

Progress through a function

00000000004004ed <loop>:

4004ed: 55

push %rbp

→ 4004ee: 48 89 e5

mov %rsp,%rbp

4004f1: c7 45 fc 00 00 00 00

movl \$0x0,-0x4(%rbp)

4004f8: 83 45 fc 01

addl \$0x1,-0x4(%rbp)

4004fc: eb fa

jmp 4004f8 <loop+0xb>

Progress through a function

00000000004004ed <loop>:

4004ed: 55

4004ee: 48 89 e5

→ 4004f1: c7 45 fc 00 00 00 00

4004f8: 83 45 fc 01

4004fc: eb fa

push %rbp

mov %rsp,%rbp

movl \$0x0,-0x4(%rbp)

addl \$0x1,-0x4(%rbp)

jmp 4004f8 <loop+0xb>

Progress through a function

00000000004004ed <loop>:

4004ed: 55

4004ee: 48 89 e5

4004f1: c7 45 fc 00 00 00 00

→ 4004f8: 83 45 fc 01

4004fc: eb fa

push %rbp

mov %rsp,%rbp

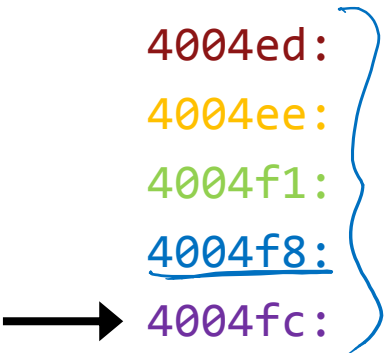
movl \$0x0,-0x4(%rbp)

addl \$0x1,-0x4(%rbp)

jmp 4004f8 <loop+0xb>

Progress through a function

0000000004004ed <loop>:



4004ed:	55
4004ee:	48 89 e5
4004f1:	c7 45 fc 00 00 00 00
<u>4004f8:</u>	83 45 fc 01
→ 4004fc:	eb fa


push	%rbp
mov	%rsp,%rbp
movl	\$0x0,-0x4(%rbp)
addl	\$0x1,-0x4(%rbp)
jmp	<u>4004f8</u> <loop+0xb>

Progress through a function

00000000004004ed <loop>:

→ 4004ed: 55 push
4004ee: 48 89 e5 mov
4004f1: c7 45 fc 00 00 00 00 movl
4004f8: 83 45 fc 01 addl
4004fc: eb fa jmp

4004fd	fa
4004fc	eb
4004fb	01
4004fa	fc
4004f9	45
4004f8	83
4004f7	00
4004f6	00
4004f5	00
4004f4	00
4004f3	fc
4004f2	45
4004f1	c7
<u>4004f0</u>	e5
<u>4004ef</u>	89
<u>4004ee</u>	48
4004ed	55



Progress through a function

00000000004004ed <loop>:


4004ed:55push

→ 4004ee:48 89 e5mov

4004f1:c7 45 fc 00 00 00 00movl

4004f8:83 45 fc 01addl

4004fc:eb fajmp



4004fd	fa
4004fc	eb
4004fb	01
4004fa	fc
4004f9	45
4004f8	83
4004f7	00
4004f6	00
4004f5	00
4004f4	00
4004f3	fc
4004f2	45
4004f1	c7
4004f0	e5
4004ef	89
4004ee	48
4004ed	55

Progress through a function

00000000004004ed <loop>:

```
4004ed:      55                push
4004ee:      48 89 e5        mov
→ 4004f1:      c7 45 fc 00 00 00 00  movl
4004f8:      83 45 fc 01      addl
4004fc:      eb fa          jmp
```



4004fd	fa
4004fc	eb
4004fb	01
4004fa	fc
4004f9	45
4004f8	83
4004f7	00
4004f6	00
4004f5	00
4004f4	00
4004f3	fc
4004f2	45
4004f1	c7
4004f0	e5
4004ef	89
4004ee	48
4004ed	55

Progress through a function

00000000004004ed <loop>:

→

4004ed:	55	
4004ee:	48 89 e5	
4004f1:	c7 45 fc 00 00 00 00	
4004f8:	83 45 fc 01	
4004fc:	eb fa	

push
mov
movl
addl
jmp



4004fd	fa
4004fc	eb
4004fb	01
4004fa	fc
4004f9	45
4004f8	83
4004f7	00
4004f6	00
4004f5	00
4004f4	00
4004f3	fc
4004f2	45
4004f1	c7
4004f0	e5
4004ef	89
4004ee	48
4004ed	55

Progress through a function

00000000004004ed <loop>:

4004ed:

55

4004ee:

48 89 e5

4004f1:

c7 45 fc 00 00 00 00

4004f8:

83 45 fc 01

4004fc:

eb fa


push

mov

movl

addl

jmp



This special register has a name:

%rip

Special hardware is responsible for setting its value to advance to the next instruction—very similar logic to the disassemble part of your assign4 that looks at opcode to determine how many subsequent bytes are there.

4004fd	fa
4004fc	eb
4004fb	01
4004fa	fc
4004f9	45
4004f8	83
4004f7	00
4004f6	00
4004f5	00
4004f4	00
4004f3	fc
4004f2	45
4004f1	c7
4004f0	e5
4004ef	89
4004ee	48
4004ed	55

“Interfering” with %rip

IF %rip ALWAYS ADVANCES TO NEXT INSTRUCTION, HOW DO WE
“SKIP” INSTRUCTIONS IN AN IF-ELSE, OR REPEAT INSTRUCTIONS
IN A LOOP?


The jmp instruction

0000000004004ed <loop>:

```
4004ed:      push    %rbp
4004ee:      mov     %rsp,%rbp
4004f1:      movl    $0x0,-0x4(%rbp)
4004f8:      addl    $0x1,-0x4(%rbp)
→ 4004fc:      jmp     4004f8 <loop+0xb>
```



%rip



4004fd	fa
4004fc	eb
4004fb	01
4004fa	fc
4004f9	45
4004f8	83
4004f7	00
4004f6	00
4004f5	00
4004f4	00
4004f3	fc
4004f2	45
4004f1	c7
4004f0	e5
4004ef	89
4004ee	48
4004ed	55


The jmp instruction

0000000004004ed <loop>:

```
4004ed:      push    %rbp
4004ee:      mov     %rsp,%rbp
4004f1:      movl    $0x0,-0x4(%rbp)
→ 4004f8:      addl    $0x1,-0x4(%rbp)
4004fc:      jmp     4004f8 <loop+0xb>
```



%rip



4004fd	fa
4004fc	eb
4004fb	01
4004fa	fc
4004f9	45
4004f8	83
4004f7	00
4004f6	00
4004f5	00
4004f4	00
4004f3	fc
4004f2	45
4004f1	c7
4004f0	e5
4004ef	89
4004ee	48
4004ed	55


The jmp instruction

0000000004004ed <loop>:

```
4004ed:      push    %rbp
4004ee:      mov     %rsp,%rbp
4004f1:      movl    $0x0,-0x4(%rbp)
4004f8:      addl    $0x1,-0x4(%rbp)
→ 4004fc:      jmp     4004f8 <loop+0xb>
```



%rip



4004fd	fa
4004fc	eb
4004fb	01
4004fa	fc
4004f9	45
4004f8	83
4004f7	00
4004f6	00
4004f5	00
4004f4	00
4004f3	fc
4004f2	45
4004f1	c7
4004f0	e5
4004ef	89
4004ee	48
4004ed	55


The jmp instruction

0000000004004ed <loop>:

```
4004ed:      push    %rbp
4004ee:      mov     %rsp,%rbp
4004f1:      movl    $0x0,-0x4(%rbp)
→ 4004f8:      addl    $0x1,-0x4(%rbp)
4004fc:      jmp     4004f8 <loop+0xb>
```



%rip



4004fd	fa
4004fc	eb
4004fb	01
4004fa	fc
4004f9	45
4004f8	83
4004f7	00
4004f6	00
4004f5	00
4004f4	00
4004f3	fc
4004f2	45
4004f1	c7
4004f0	e5
4004ef	89
4004ee	48
4004ed	55


The jmp instruction

0000000004004ed <loop>:

```
4004ed:      push    %rbp
4004ee:      mov     %rsp,%rbp
4004f1:      movl    $0x0,-0x4(%rbp)
4004f8:      addl    $0x1,-0x4(%rbp)
→ 4004fc:      jmp     4004f8 <loop+0xb>
```



%rip



4004fd	fa
4004fc	eb
4004fb	01
4004fa	fc
4004f9	45
4004f8	83
4004f7	00
4004f6	00
4004f5	00
4004f4	00
4004f3	fc
4004f2	45
4004f1	c7
4004f0	e5
4004ef	89
4004ee	48
4004ed	55


The jmp instruction

0000000004004ed <loop>:

```
4004ed:      push    %rbp
4004ee:      mov     %rsp,%rbp
4004f1:      movl    $0x0,-0x4(%rbp)
→ 4004f8:      addl    $0x1,-0x4(%rbp)
4004fc:      jmp     4004f8 <loop+0xb>
```



%rip



4004fd	fa
4004fc	eb
4004fb	01
4004fa	fc
4004f9	45
4004f8	83
4004f7	00
4004f6	00
4004f5	00
4004f4	00
4004f3	fc
4004f2	45
4004f1	c7
4004f0	e5
4004ef	89
4004ee	48
4004ed	55


The jmp instruction

0000000004004ed <loop>:

```
4004ed:      push    %rbp
4004ee:      mov     %rsp,%rbp
4004f1:      movl    $0x0,-0x4(%rbp)
4004f8:      addl    $0x1,-0x4(%rbp)
→ 4004fc:      jmp     4004f8 <loop+0xb>
```



%rip



4004fd	fa
4004fc	eb
4004fb	01
4004fa	fc
4004f9	45
4004f8	83
4004f7	00
4004f6	00
4004f5	00
4004f4	00
4004f3	fc
4004f2	45
4004f1	c7
4004f0	e5
4004ef	89
4004ee	48
4004ed	55


The jmp instruction

0000000004004ed <loop>:

```
4004ed:      push    %rbp
4004ee:      mov     %rsp,%rbp
4004f1:      movl    $0x0,-0x4(%rbp)
→ 4004f8:      addl    $0x1,-0x4(%rbp)
4004fc:      jmp     4004f8 <loop+0xb>
```



%rip

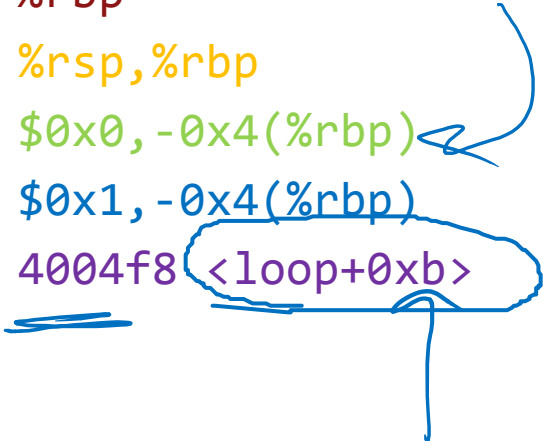


4004fd	fa
4004fc	eb
4004fb	01
4004fa	fc
4004f9	45
4004f8	83
4004f7	00
4004f6	00
4004f5	00
4004f4	00
4004f3	fc
4004f2	45
4004f1	c7
4004f0	e5
4004ef	89
4004ee	48
4004ed	55


C code for our example

0000000004004ed <loop>:

```
4004ed:    push    %rbp
4004ee:    mov     %rsp,%rbp
4004f1:    movl    $0x0,-0x4(%rbp)
4004f8:    addl    $0x1,-0x4(%rbp)
4004fc:    jmp     4004f8 <loop+0xb>
```



```
void loop()
{
    int i = 0;
again:
    ++i;
    goto again;
}
```



Conditional jumps

Typical 2-step control flow

1. Compare two values to **write** the condition codes (implicit destination register)
 - › `cmp, test`
2. Conditionally jump based on **reading** the condition codes (implicit source register)
 - › `je, jne, jl, jg`



- There is also a 1-step unconditional jump
- Doesn't look at condition code, just goes no matter what
 - › `jmp [target]`

STEP 1 of control flow: cmp, test

$$107 - 6 = 101$$

Op	Source1	Source2	Dest Comments
cmp	<u>op2</u>	<u>op1</u>	op1 – op2, sets condition codes
test	op2	op1	op1 & op2, sets condition codes

- op1 and op2 can be any of the complex addressing modes we've seen
- Implicit destination %eflags contains condition codes
 - › Sequence of Boolean values packed into one register
 - › **t** is the result of the cmp or test operation above
 - ZF = zero flag (t = 0)
 - SF = sign flag (t < 0)
 - CF = carry flag (there was a carry out of MSB*, *i.e.* unsigned overflow)
 - OF = overflow flag (MSB* changed from 0 to 1, *i.e.* signed overflow)
 - ...

* MSB = "Most Significant Bit"

What is the value of %eflags after this code?

```
00000000004004fe <if_then>:
4004fe:      push    %rbp
4004ff:      mov     %rsp,%rbp
400502:      mov     %edi,0x4(%rbp)
400505:      cmpl     $0x6,-0x4(%rbp)
...
```

Which of these bits (flags) are 1 (set) if we pass argument 107 to this function?

- ZF = zero flag ($t = 0$)
- SF = sign flag ($t < 0$)
- CF = carry flag (there was a carry out of MSB*, *i.e.* unsigned overflow)
- OF = overflow flag (MSB* changed from 0 to 1, *i.e.* signed overflow)

STEP 2 of control flow: jump

Op	Target	Remarks
jmp	target	Unconditional jump
je	target	Jump if ZF is 1, in other words $op1 - op2 = 0$ in previous cmp, in other words $op1 = op2$

- Target is a memory address: the address of the instruction where we should jump
- Implicit source %eflags contains condition codes
 - › Sequence of Boolean values packed into one register
 - ZF = zero flag
 - SF = sign flag
 - CF = carry flag
 - OF = overflow flag
 - ...



Control operations

<code>cmpl op2, op1</code>	# result = op1 - op2, discards result, sets condition
<code>test op2, op1</code>	# result = op1 & op2, discards result, sets condition
<code>jmp target</code>	# unconditional jump
<code>je target</code>	# jump equal, synonym jz jump zero (ZF=1)
<code>jne target</code>	# jump not equal, synonym jnz (ZF=0)
<code>jnl target</code>	# jump less than, synonym jnge (SF!=OF)
<code>jle target</code>	# jump less or equal, synonym jng (ZF=1 or SF!=OF)
<code>jg target</code>	# jump greater than, synonym jnle (ZF=0 and SF=OF)
<code>jge target</code>	# jump greater or equal, synonym jnl (SF=OF)
<code>ja target</code>	# jump above, synonym jnbe (CF=0 and ZF=0)
<code>jb target</code>	# jump below, synonym jnae (CF=1)
<code>js target</code>	# jump signed (SF=1)
<code>jns target</code>	# jump not signed (SF=0)

What is the value of %rip after the jne?

0000000004004fe <if_then>:

```
4004fe:      push    %rbp
4004ff:      mov     %rsp,%rbp
400502:      mov     %edi,-0x4(%rbp)
400505:      cmpl    $0x6,-0x4(%rbp)
400509:      jne     40050f
40050b:      addl    $0x1,-0x4(%rbp)
```

...

What is param1 at the marked location when the input was 3?

- a) 400509
- b) 40050b
- c) 40050f
- d) Something else

Target instruction

Op	Target	Remarks
jmp	target	Unconditional jump
je	target	Jump if ZF is 1, in other words <code>op1-op2=0</code> in previous <code>cmp</code> , in other words <code>op1=op2</code>

- Reminder: everything is bits/bytes to a computer!
 - › Instructions are just 1s and 0s that we interpret in a certain way
 - › Those bits/bytes are in a memory location, just like pointers, ints, floats, doge pictures, cat videos, and other data
 - › So “target” is an address (or offset from current address) to write to the PC (program counter)
 - char = 1 byte
 - int = 4 bytes
 - Double, void* = 8 bytes
 - instruction = ?? bytes # **LET'S LOOK IN GDB TO FIND OUT**
- Stanford University

Instruction Set Architectures

SOME CONTEXT AND TERMINOLOGY

Instruction Set Architecture

- **The ISA defines:**

- › Operations that the processor can execute
- › Data transfer operations + how to access data
- › Control mechanisms like branch, jump (think loops and if-else)
- › Contract between programmer/compiler and hardware

- **Layer of abstraction:**

- › Above:
 - Programmer/compiler can write code for the ISA
 - New programming languages can be built on top of the ISA as long as the compiler will do the translation
- › Below:
 - New hardware can implement the ISA
 - Can have even potentially radical changes in hardware implementation
 - Have to “do” the same thing from programmer point of view


- **ISAs have incredible inertia!**

- › Legacy support is a huge issue for x86-64

Two major categories of Instruction Set Architectures

■ CISC:

- › **Complex** instruction set computers
 - e.g., x86 (**CS107 studies this**)
- › Have special instructions for each thing you might want to do
- › Can write code with fewer instructions, because each instruction is very expressive



■ RISC:

- › **Reduced** instruction set computers
 - e.g., MIPS
- › Have only a very tiny number of instructions, optimize the heck out of them in the hardware
- › Code may need to be longer because you have to go roundabout ways of achieving what you wanted



C translation examples

If statement construction

```
/* if-stmt */
void if_then(int param1)
{
    if (param1 == 6)
        param1++;
    param1 *= 2;
}

# gcc output
0000000004004fe <if_then>:
    4004fe:      push    %rbp
    4004ff:      mov     %rsp,%rbp
    400502:      mov     %edi,-0x4(%rbp)
    400505:      cmpl    $0x6,-0x4(%rbp)
    400509:      jne     40050f
    40050b:      addl    $0x1,-0x4(%rbp)
    40050f:      shll    -0x4(%rbp)
    400512:      pop     %rbp
    400513:      retq
```

If-else construction: Find the bug!

/* if-else */	# BUGGY (not actual gcc output, but close)
void if_else(int param1)	0000000000400514 <if_else>:
{	400514: push %rbp
if (param1 < 5)	400515: mov %rsp,%rbp
param1++;	400518: mov %edi,-0x4(%rbp)
else	40051b: cmpl \$0x4,-0x4(%rbp)
param1--;	40051f: jg 400527
param1 = -param1;	400521: addl \$0x1,-0x4(%rbp)
//what is param1 here?	400527: subl \$0x1,-0x4(%rbp)
}	40052b: negl -0x4(%rbp)
	40052e: pop %rbp
	40052f: retq

What is param1 at the marked location when the input was 3?

- (a) 2 (b) 3 (c) 4 (d) something else

If-else construction

```
/* if-else */  
if (num > 3) {  
    x = 10;  
} else {  
    x = 7;  
}
```

```
/* equivalent if-else */  
if (num <= 3) GOTO L2  
x = 10;  
GOTO L3  
L2: x = 7;  
L3: ...
```

```
# equivalent AMD64 pseudocode  
Test  
Branch OVER if-body if test fails  
If-body  
jmp past else-body  
Else-body  
[PAST ELSE-BODY]
```

For loop construction

```
/* for loop */  
for (int i=0; i<n; i++) {  
    /* body */  
}
```

```
/* equivalent while loop */  
int i=0;  
while (i<n) {  
    /* body */  
    i++;  
}
```

```
# equivalent AMD64 pseudocode  
Initialization  
Test  
Branch past loop if fails  
Body  
Increment  
jmp to Test
```

For loop construction

equivalent AMD64

Initialization

Test

Branch past loop if fails

Body

Increment

jmp to Test

AMD64 gcc actually emits

Initialization

jmp to Test

Body

Increment

Test

Branch to Body if succeeds

- Same number of instructions! Why does gcc use the format on the right?

Say for loop “for (int i=0; i<n; i++)” and n=0, n=1000

Compare the instructions executed in the left and right

- A. LEFT and RIGHT have same number of instructions
- B. LEFT has more instructions (bad for left)
- C. RIGHT has more instructions (bad for right)
- D. Other/help

Computer Architecture BIG IDEA:

Code with Smaller Static Instruction Count
!= Code with Smaller Dynamic Instruction Count

- Our two codes had the same number of instructions
 - › **Same** static instruction count
- If loop never executes, right had **higher** dynamic instruction count (bad for right)
- If loop executes many times, left had **higher** dynamic instruction count (bad for left)
- **This lack of correlation is very common!**
 - › There are even cases where the compiler emits a static instruction count that is *several times* longer than an alternative, yet still more efficient assuming loops execute many times (e.g. loop unrolling)

Discussion question:

- Does the compiler **know** that the loop will execute many times?
 - › In general, no!
- So...what if our code has loops that always execute a small number of times? Did gcc make a bad decision?