

# Computer Systems

CS107

Cynthia Lee

# Today's Topics

- › Managing the heap
- › Important for your final assignment, assign7!



# Heap Manager

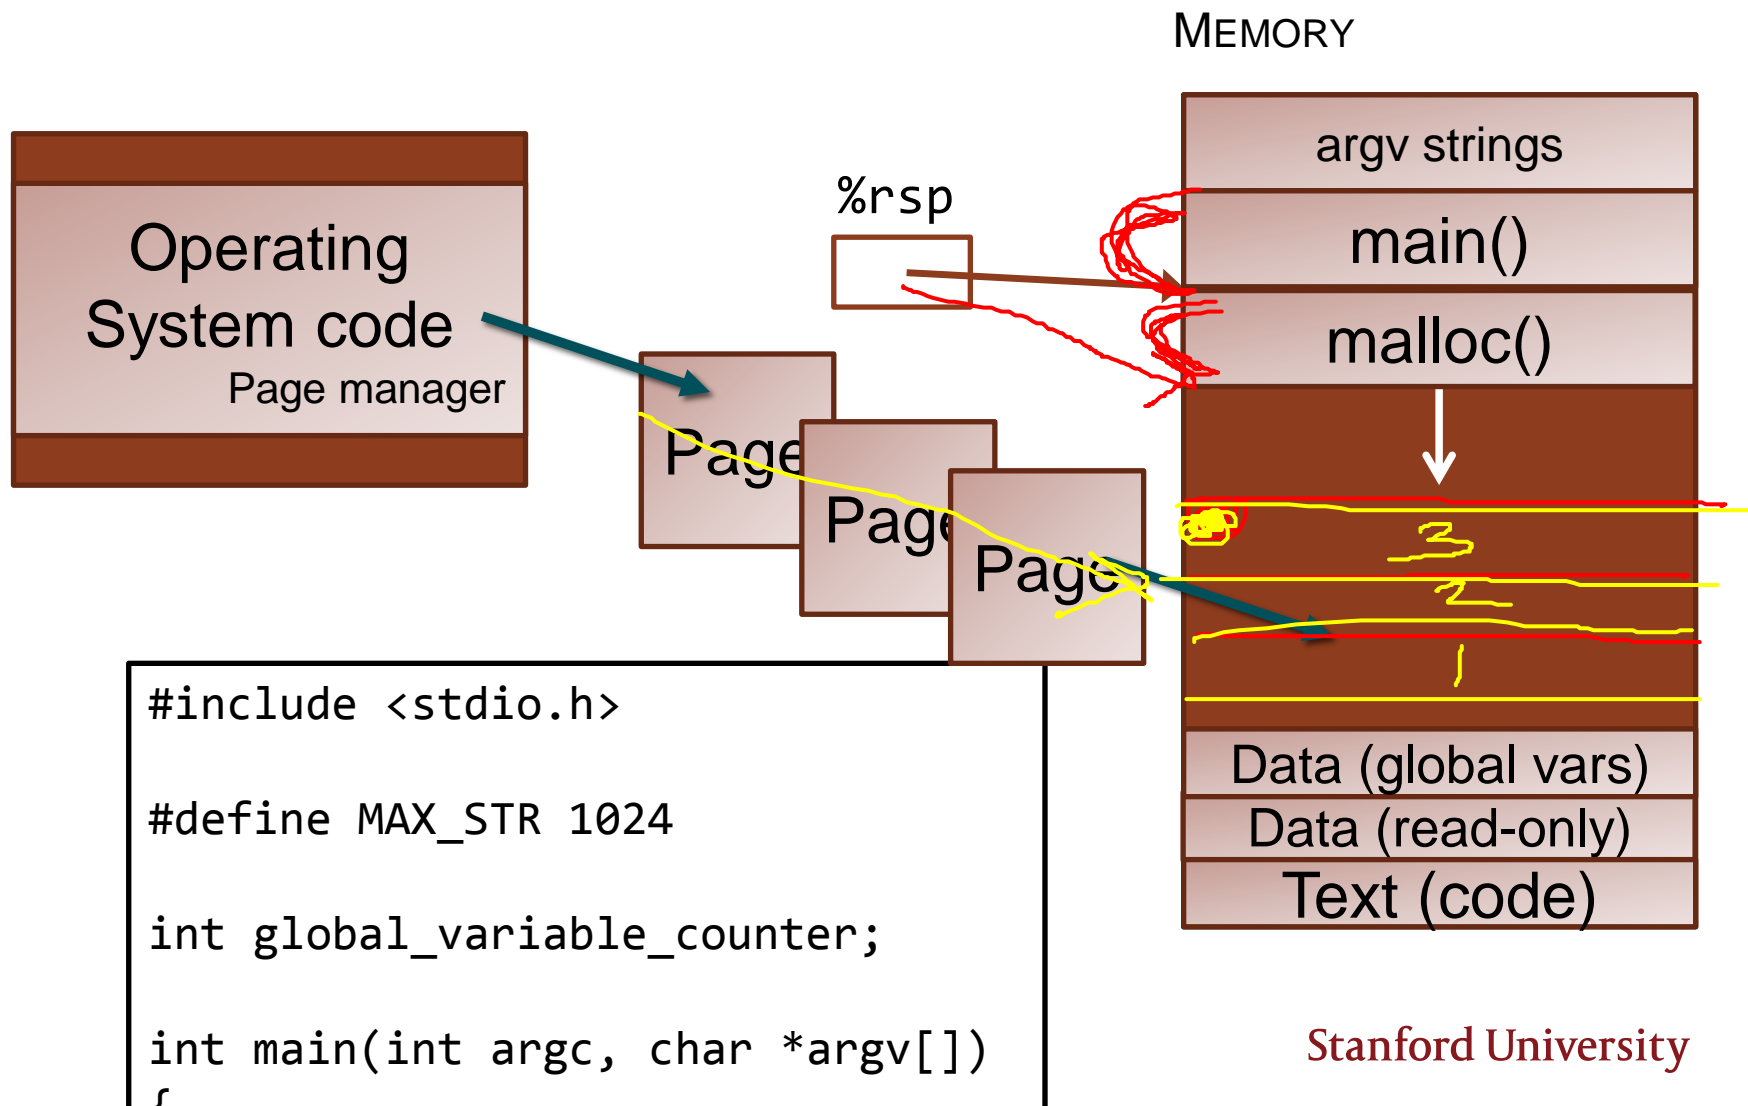
Remember what you know from being a client of malloc/free:

The heap manager functions like a hotel front desk—handing out room assignments based on size needs (malloc), and checking people out when they're ready to leave (free).

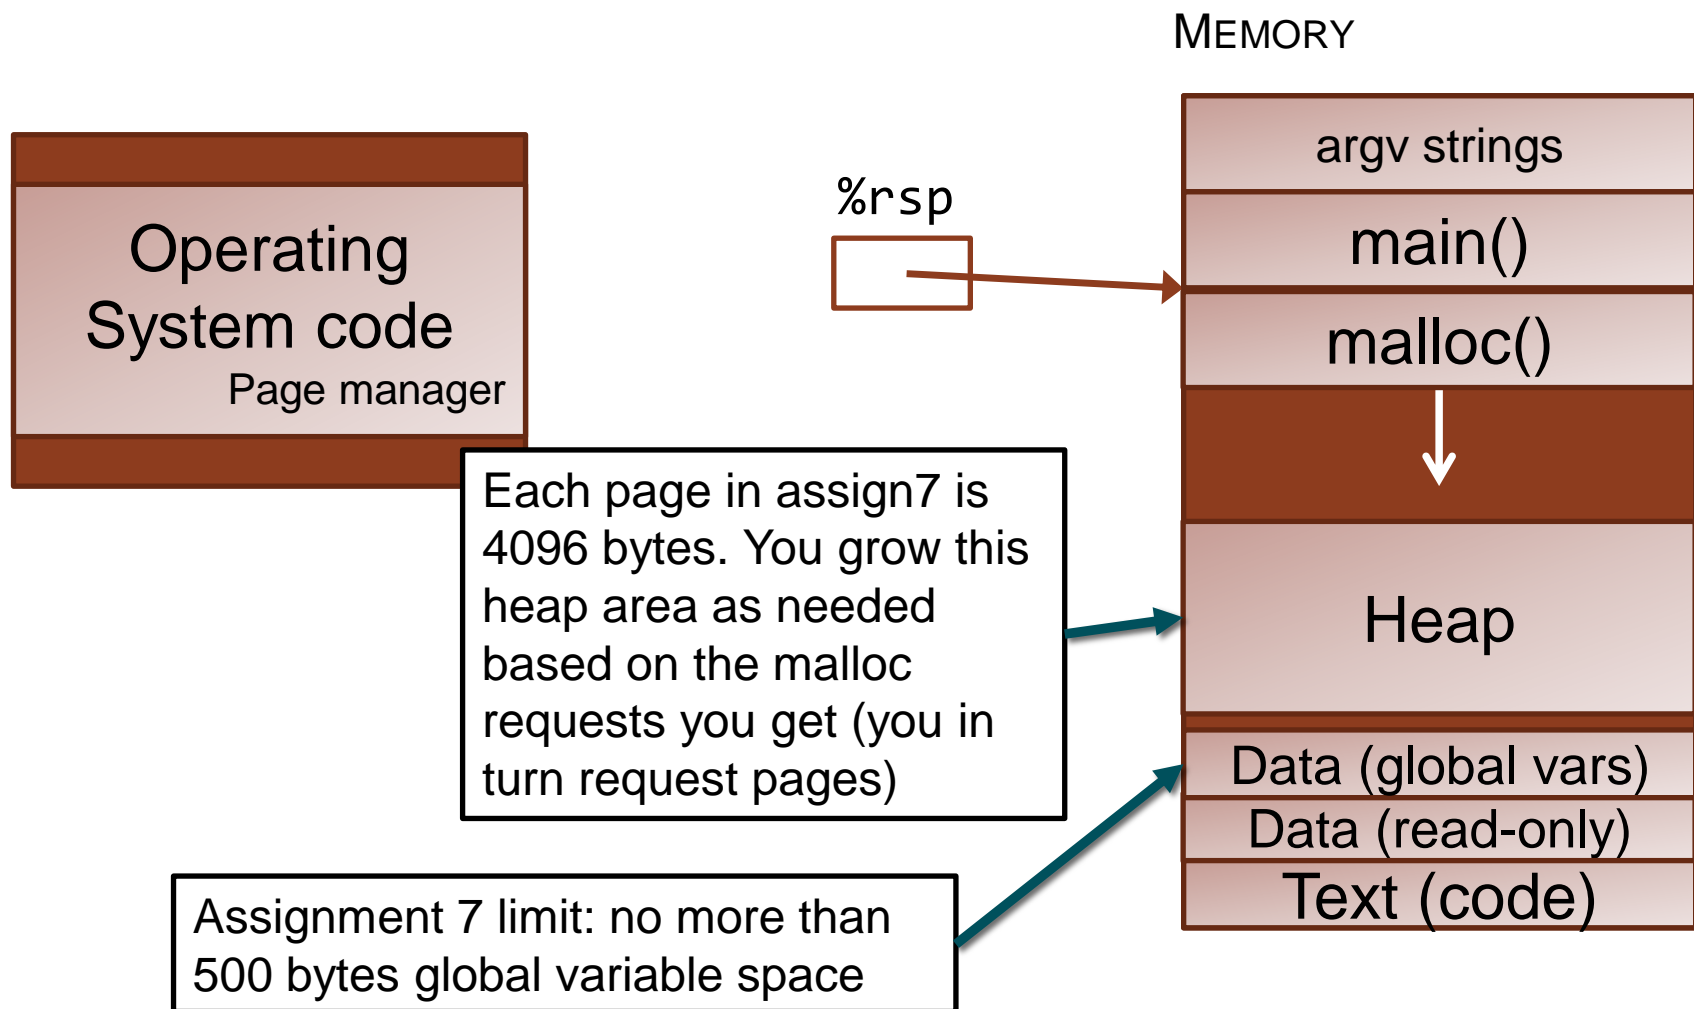
## Heap Allocator: you (almost) don't need a spec for this assignment!

- You will implement three functions, whose arguments, return values, and behavior, you already know:
  - `void *malloc(size_t num_bytes_requested)`
  - `void *realloc(void *old_ptr, size_t new_bytes_requested)`
  - `void free(void *ptr)`
- Other requirements/limits:
  - Maximum request size for a single request is INT\_MAX bytes
  - Must be fast (throughput) and make efficient use of space (utilization)
  - No more than 500 bytes of global variables, other storage comes out of the heap memory you would otherwise give to user.

## Big picture: what do you have to work with?



## Big picture: what do you have to work with?





# Heap Manager

Some simple implementation framework examples

## Heap Manager: Example client code

```
void *a, *b, *c, *d, *e, *f;  
a = malloc(4);  
b = malloc(8);  
c = malloc(4);  
d = malloc(4);  
free(a);  
free(c);  
e = malloc(12);  
b = realloc(b, 12);  
d = realloc(d, 8);  
f = malloc(12);
```



## Implementation #1: global variable space lists

```
void *a, *b, *c, *d, *e, *f;
```

```
a = malloc(4);
```

```
b = malloc(8);
```

```
c = malloc(4);
```

```
d = malloc(4);
```

```
free(a);
```


```
free(c);
```

```
e = malloc(12);
```

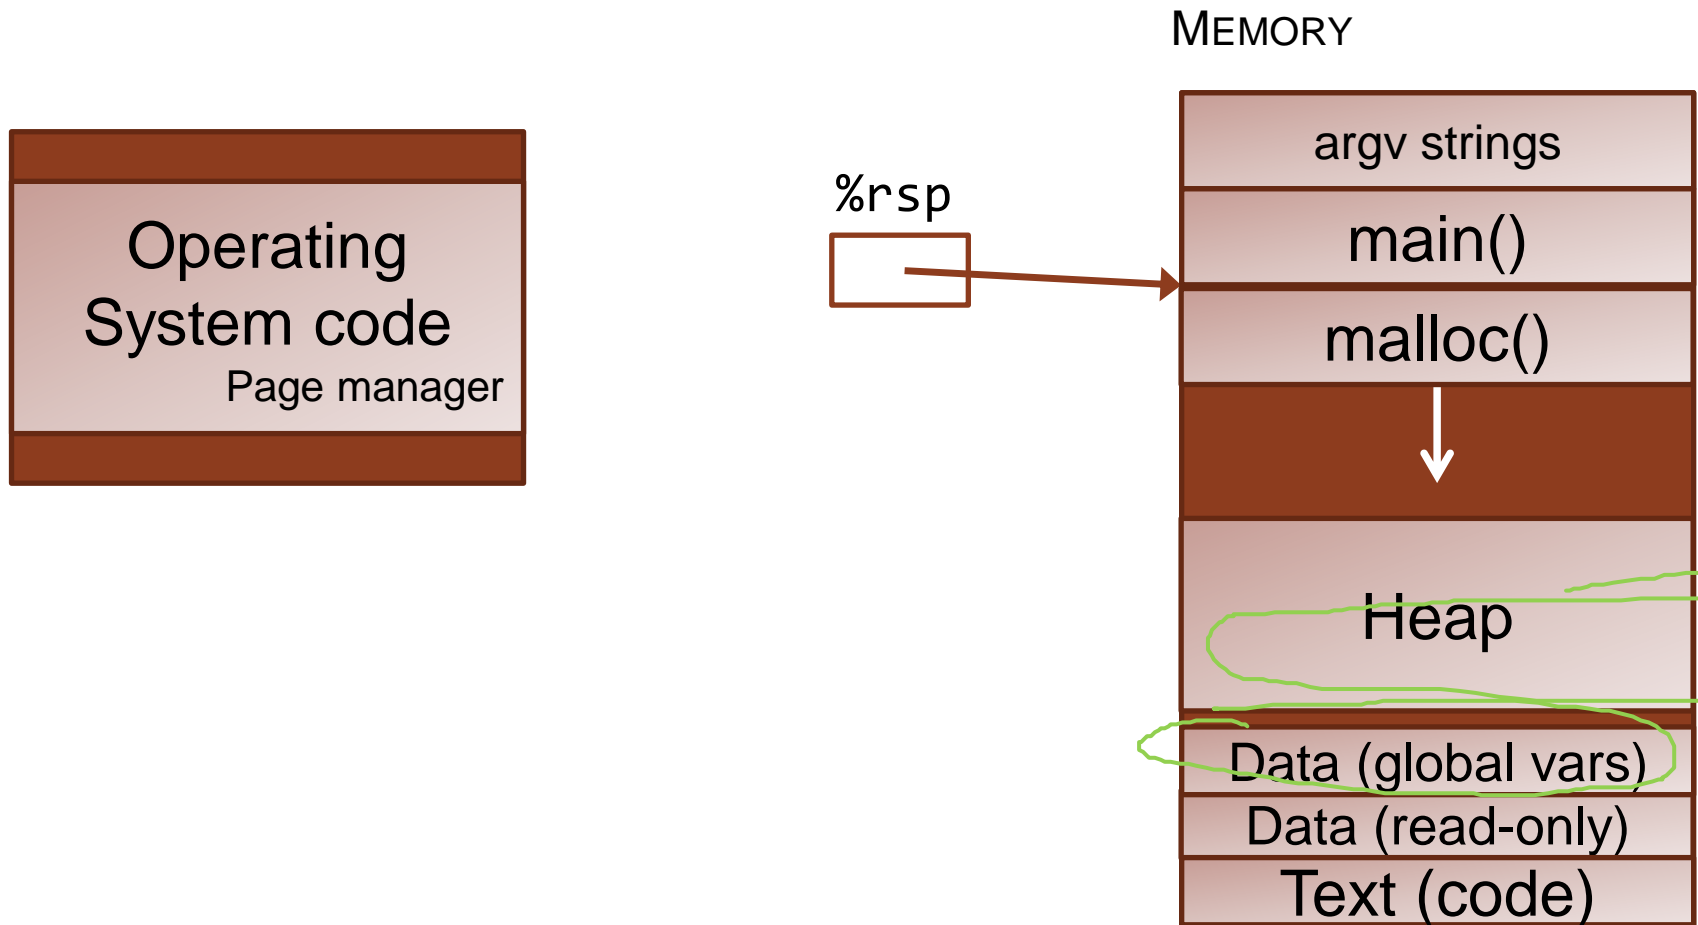
```
b = realloc(504b, 12);
```

```
d = realloc(d, 8);510
```

```
f = malloc(12);
```



## Big picture: what do you have to work with?



## Discussion question

### Could the heap manager pause and do a “defrag” operation on the heap at this point?

- › Move data in allocated areas over to coalesce free space into a contiguous block
- › Don't change size of any allocated block, and carefully copy data

A. YES, great idea!

B. YES it can be done, but not a good idea for some reason (e.g., not efficient use of time)

C. NO, it can't be done!

Why or why not?

## Implementation #2: Basic pre-node headers

```
void *a, *b, *c, *d, *e, *f;
```

```
a = malloc(4);
```

```
b = malloc(8);
```

```
c = malloc(4);
```

```
d = malloc(4);
```

```
free(a);
```

```
free(c);
```

```
e = malloc(12);
```

```
b = realloc(b, 12);
```

```
d = realloc(d, 8);
```

```
f = malloc(12);
```

Discussion question

How do we handle `realloc(b, 12)`?

53c

538

534

530

52c

528

524

520

51c

518

514

510

50c

508

504

0x500

32 + N

24 + Y

8 + N

4 + Y

## Discussion question

**Is there anything about these approaches to the heap manager (generally, or specifically this pre-node header implementation) that strikes you as being a bit dangerous?**

- i.e., from a security and software engineering perspective



What kind of accidental (or malicious?) bugs in the user's code could impact the heap and heap manager? How?

What defenses does the heap manager have available?



## Implementation #3: Pre-node headers + free nodes organized into a linked list ("free list")

```
void *a, *b, *c, *d, *e, *f;
```

```
a = malloc(4);
```

```
b = malloc(8);
```

```
c = malloc(4);
```

```
d = malloc(4);
```

```
free(a);
```

```
free(c);
```

```
e = malloc(12);
```

```
b = realloc(b, 12);
```

```
d = realloc(d, 8);
```

```
f = malloc(12);
```

53c

538

534

530

52c

528

524

520

51c

518

514

510

50c

508

504

0x500

# Free List

Global variable memory:


free\_list:



beg :

0x500

Heap memory (in pages):





# Heap Manager

Optimizations



## Back to the “defrag” idea:

- › We know that we can’t rearrange pointers at will for currently-allocated memory
- › However, there is nothing to stop us from taking adjacent already-freed blocks and coalescing them into a larger freed block

YES, great idea!

The question is, given our free list is a linked list, how do we “notify” incoming pointers to free blocks that the block has now been aggregated into one conglomerate node? Hm....

- See textbook’s “footer” solution


# Free List

Global variable memory:

`free_list:`




Heap memory (in pages):



## Optimisation: Explicit free list *bucketed by size*

Global variable memory:

Heap memory (in pages):



How can we balance the conflicting goals of the heap manager?

**Correctness**

(obviously)

**Throughput**


(speed of handling requests)




**Utilization**

(not too many holes in space)


# Your Pinterest Board of Design Ideas for Heap Allocator




**DIY Easy!**  
In-Use List +  
Free List



**Cute Storage Idea!**  
Header holds  
block metadata  
adjacent to  
payload



**Declutter Tip**  
Coalesce  
adjacent free  
blocks to reduce  
fragmentation



**Organize!**  
Multiple free  
lists, bucketed  
by block size