

Computer Systems

CS107

Cynthia Lee

Today's Topics

TODAY'S LECTURE:

- › Performance optimization and the memory hierarchy

ANNOUNCEMENTS:

- › Assign6 & 7 NO late days allowed!
 - **Submit early and often!**
 - Submit whatever you have that could be worth any points at all as soon as you have it, re-submit any time you improve it
 - That way if a catastrophe happens at deadline time and you can't submit, at least you have some points from your earlier submissions!

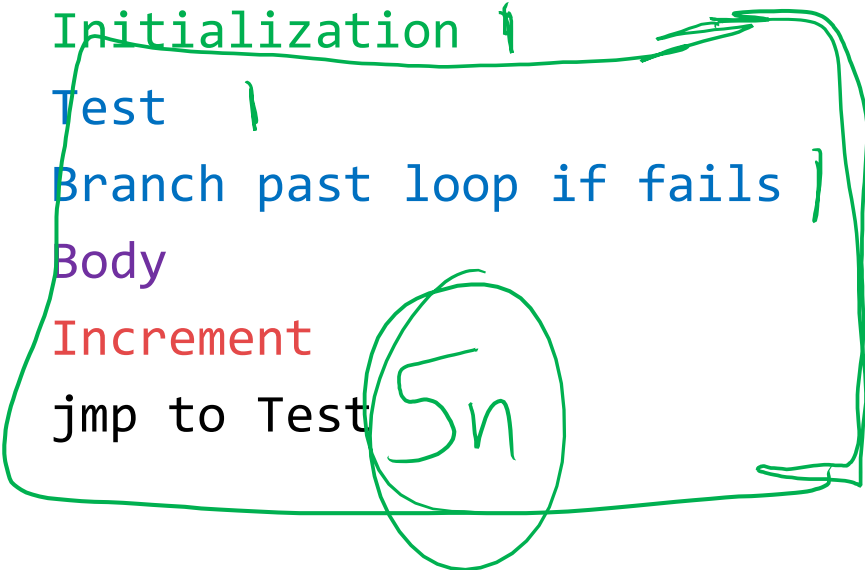
Valgrind for instruction profiling

Your code is awaited at the gates of Valgrindhalla! It shall ride eternal, shiny and chrome.

For loop construction

straightforward assembly

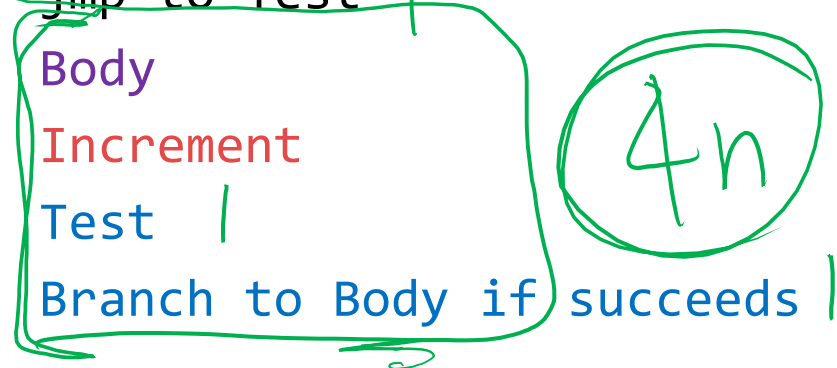
```
Initialization |  
Test |  
Branch past loop if fails |  
Body  
Increment  
jmp to Test
```



$5n$

what gcc actually emits

```
Initialization |  
jmp to Test |  
Body  
Increment  
Test |  
Branch to Body if succeeds |
```



$4n$

Automated tool for creating dynamic instruction counts

```
/* Simple selection sort algorithm,  $O(N^2)$  */
void selsort(int *arr, int n)
{
    4,005      for (int i = 0; i < n; i++) {
    2,000          int min = i;
2,005,000      for(int j = i+1; j < n; j++)
7,003,974          if (arr[j] < arr[min]) min = j;
    13,000      swap(&arr[i], &arr[min]);
                }
    }
```

Valgrind can show you dynamic instruction counts

- Valgrind is a tool for memory profiling!
- But wait, there's more! Can also do other kinds of performance profiling

```
% valgrind --tool=callgrind --simulate-cache=yes ./array
```

```
% callgrind_annotate --auto=yes callgrind.out.[procID]
```

- › The last argument in the first call is the name of the executable to profile (*array is only an example*)
- › The last argument in the second call is the name of the output file created by the first call (procID is a number, will be different each time)
- › Add `--inclusive=yes` if you want to include the cost of function calls (by default it only counts code actually in the current function)

Performance example: sorting

Cycle counters can be a more precise way of capturing intuitions similar to what Big-O analysis captures

Sorting algorithms reminder

```
/* Simple selection sort algorithm,  $O(N^2)$  */  
void selsort(int *arr, int n)
```

```
{  
    for (int i = 0; i < n; i++) {  
        int min = i;  
        for(int j = i+1; j < n; j++)  
            if (arr[j] < arr[min]) min = j;  
        swap(&arr[i], &arr[min]);  
    }  
}
```

```
/* Simple insertion sort algorithm,  $O(N^2)$  */  
void inssort(int *arr, int n)
```

```
{  
    for (int i = 1; i < n; i++) {  
        for (int j = i; j > 0 && arr[j-1] > arr[j]; j--) {  
            swap(&arr[j], &arr[j-1]);  
        }  
    }  
}
```


Sorting algorithms reminder: Big-O

- Quicksort: $O(n \log n)$
 - › Assuming not consistently bad pivot selection
- Selection Sort: $O(n^2)$
- Insertion Sort: $O(n^2)$
- We can empirically verify our CS106B knowledge of these algorithms by using cycle counter for a few array sizes:

```
myth> ./sorts
Sorting 5000 elements:
qsort cycles 2.78M (opt)
selsort cycles 26.66M (opt)
```

```
Sorting 10000 elements:
qsort cycles 4.04M (opt)
selsort cycles 105.15M (opt)
```

```
Sorting 20000 elements:
qsort cycles 8.53M (opt)
selsort cycles 429.05M (opt)
```

Interpreting cycle counts

SORTING 1000 ELEMENTS:

- Alg-A 0.88M cycles
 - Alg-B 5.31M cycles
 - Alg-B 5.07M cycles (on already-sorted data)
 - Alg-C 7.14M cycles
 - Alg-C 0.01M cycles (on already-sorted data)
-
- Guess the algorithm:
 - A. A = quicksort, B = selection sort, C = insertion sort
 - B. A = quicksort, B = insertion sort, C = selection sort
 - C. Something else

Performance example: array access

We'll look at a few different patterns of array access

We have code that accesses an array in different ways

$O(n)$ for (int i = 0; i < n; i++) //normal access
sum += a[i];

$O(n)$ for (int i = n-1; i >= 0; i--) //backwards access
sum += a[i];

$O(n)$ for (int i = 0; i < n; i+=2) //every other one (evens)
sum += a[i];
for (int i = 1; i < n; i+=2) //then odds
sum += a[i];

$O(n)$ for (int i = 0; i < n; i++) //random order
sum += a[indexes[i]];

You can probably guess that the punchline is that these have very different performance profiles (they do).

We have code that accesses an array in different ways

1 for (int i = 0; i < n; i++) //forwards
 sum += a[i];

2 for (int i = n-1; i >= 0; i--) //backwards
 sum += a[i];

3 for (int i = 0; i < n; i+=2) //evens/odds
 sum += a[i];
for (int i = 1; i < n; i+=2)
 sum += a[i];

4 for (int i = 0; i < n; i++) //random order
 sum += a[indexes[i]];

How will their cycle counts rank? (guess)

- A. $1 < 2 < 3 < 4$
- B. $1 = 2 < 3 < 4$
- C. $1 < 2 = 3 < 4$
- D. $1 = 2 = 3 < 4$
- E. Something else

Taking a look at the assembly: FORWARD

080485f0 <sum_forward>:

80485f0:	53	push	%ebx
80485f1:	b9 00 00 00 00	mov	\$0x0,%ecx
80485f6:	bb 00 00 00 00	mov	\$0x0,%ebx
80485fb:	eb 06	jmp	8048603 <sum_forward+0x13>
80485fd:	03 1c 88	add	(%eax,%ecx,4),%ebx
8048600:	83 c1 01	add	\$0x1,%ecx
8048603:	39 d1	cmp	%edx,%ecx
8048605:	7c f6	j1	80485fd <sum_forward+0xd>
8048607:	89 d8	mov	%ebx,%eax
8048609:	5b	pop	%ebx
804860a:	c3	ret	

Taking a look at the assembly: BACKWARD

0804860b <sum_backward>:

804860b:	83 ea 01	sub	\$0x1,%edx
804860e:	b9 00 00 00 00	mov	\$0x0,%ecx
8048613:	eb 06	jmp	804861b <sum_backward+0x10>
8048615:	03 0c 90	add	(%eax,%edx,4),%ecx
8048618:	83 ea 01	sub	\$0x1,%edx
804861b:	85 d2	test	%edx,%edx
804861d:	79 f6	jns	8048615 <sum_backward+0xa>
804861f:	89 c8	mov	%ecx,%eax
8048621:	c3	ret	

Taking a look at the assembly: EVEN-ODD

08048656 <sum_evenodd>:

```
8048656: 53
8048657: bb 00 00 00 00
804865c: b9 00 00 00 00
8048661: eb 06
8048663: 03 0c 98
8048666: 83 c3 02
8048669: 39 d3
804866b: 7c f6
804866d: bb 01 00 00 00
8048672: eb 06
8048674: 03 0c 98
8048677: 83 c3 02
804867a: 39 d3
804867c: 7c f6
804867e: 89 c8
8048680: 5b
8048681: c3
```

```
push    %ebx
mov     $0x0,%ebx
mov     $0x0,%ecx
jmp     8048669 <sum_evenodd+0x13>
add     (%eax,%ebx,4),%ecx
add     $0x2,%ebx
cmp     %edx,%ebx
jl      8048663 <sum_evenodd+0xd>
mov     $0x1,%ebx
jmp     804867a <sum_evenodd+0x24>
add     (%eax,%ebx,4),%ecx
add     $0x2,%ebx
cmp     %edx,%ebx
jl      8048674 <sum_evenodd+0x1e>
mov     %ecx,%eax
pop     %ebx
ret
```


Taking a look at the assembly: RANDOM

08048682 <sum_random>:

```
8048682: 57
8048683: 56
8048684: 53
8048685: bb 00 00 00 00
804868a: be 00 00 00 00
804868f: eb 09
8048691: 8b 3c 99
8048694: 03 34 b8
8048697: 83 c3 01
804869a: 39 d3
804869c: 7c f3
804869e: 89 f0
80486a0: 5b
80486a1: 5e
80486a2: 5f
80486a3: c3
```

```
push    %edi
push    %esi
push    %ebx
mov     $0x0,%ebx
mov     $0x0,%esi
jmp     804869a <sum_random+0x18>
mov     (%ecx,%ebx,4),%edi
add     (%eax,%edi,4),%esi
add     $0x1,%ebx
cmp     %edx,%ebx
jl      8048691 <sum_random+0xf>
mov     %esi,%eax
pop     %ebx
pop     %esi
pop     %edi
ret
```

And the winner is.....

```
myth> ./array
```

This program sums a 1000000-elem array. It times the traversal forward, backward, even/odd, vs randomly.

```
Sum array forward:    9.40M cycles
```

```
Sum array backward:   9.63M cycles
```

```
Sum array even/odd:   9.64M cycles
```

```
Sum array random:    36.29M cycles
```

Looking for answers in the code

```
for (int i = 0; i < n; i++) //forwards  
    sum += a[i];
```

```
for (int i = n-1; i >= 0; i--) //backwards  
    sum += a[i];
```

```
for (int i = 0; i < n; i+=2) //evens/odds  
    sum += a[i];  
for (int i = 1; i < n; i+=2)  
    sum += a[i];
```

```
for (int i = 0; i < n; i++) //random order  
    sum += a[indexes[i]];
```

forward:	9.40M cycles
backward:	9.63M cycles
even/odd:	9.64M cycles
random:	36.29M cycles

**Does anything about
this comparison
strike you as
particularly unfair?**

Making the code more fair (“apples to apples” comparison)

```
for (int i = 0; i < n; i++) //forwards
    sum += a[i];
```

```
for (int i = 0; i < n; i++) //random order
    sum += a[indexes[i]]; // must access memory TWICE!
```

```
for (int i = 0; i < n; i++) //forwards, but with 2-step index
    sum += a[indexes[i]]; // must still access TWICE
```

Taking a look at the assembly: RANDOM

08048682 <sum_random>:

8048682:	57	push	%edi
8048683:	56	push	%esi
8048684:	53	push	%ebx
8048685:	bb 00 00 00 00	mov	\$0x0,%ebx
804868a:	be 00 00 00 00	mov	\$0x0,%esi
804868f:	eb 09	jmp	804869a <sum_random+0x18>
8048691:	8b 3c 99	mov	(%ecx,%ebx,4),%edi
8048694:	03 34 b8	add	(%eax,%edi,4),%esi
8048697:	83 c3 01	add	\$0x1,%ebx
804869a:	39 d3	cmp	%edx,%ebx
804869c:	7c f3	j1	8048691 <sum_random+0xf>
804869e:	89 f0	mov	%esi,%eax
80486a0:	5b	pop	%ebx
80486a1:	5e	pop	%esi
80486a2:	5f	pop	%edi
80486a3:	c3	ret	

Taking a look at the assembly: FWD-INDEX

080486a4 <sum_fwd_ind>:

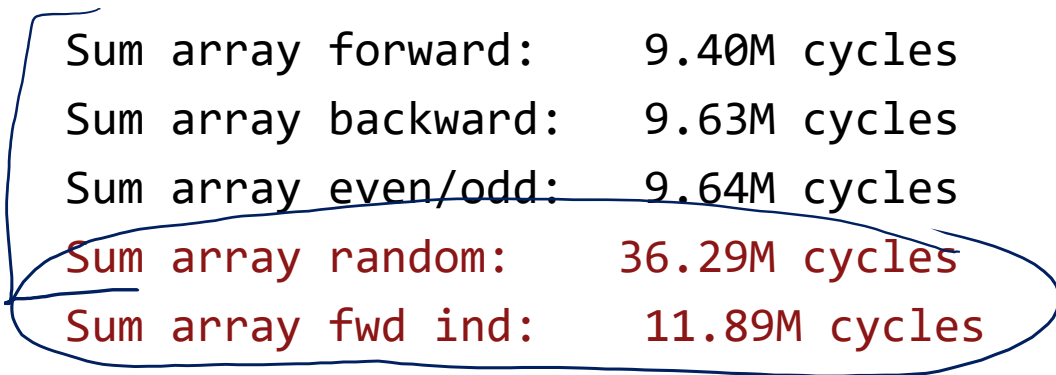
80486a4:	57	push	%edi
80486a5:	56	push	%esi
80486a6:	53	push	%ebx
80486a7:	bb 00 00 00 00	mov	\$0x0,%ebx
80486ac:	be 00 00 00 00	mov	\$0x0,%esi
80486b1:	eb 09	jmp	80486bc <sum_fwd_ind+0x18>
80486b3:	8b 3c 99	mov	(%ecx,%ebx,4),%edi
80486b6:	03 34 b8	add	(%eax,%edi,4),%esi
80486b9:	83 c3 01	add	\$0x1,%ebx
80486bc:	39 d3	cmp	%edx,%ebx
80486be:	7c f3	j1	80486b3 <sum_fwd_ind+0xf>
80486c0:	89 f0	mov	%esi,%eax
80486c2:	5b	pop	%ebx
80486c3:	5e	pop	%esi
80486c4:	5f	pop	%edi
80486c5:	c3	ret	

100 cycles

We have code that accesses an array in different ways

```
myth10:/usr/class/cs107/samples/lect16> ./array
```

This program sums a 1000000-elem array. It times the traversal forward, backward, even/odd, vs randomly.




Sum array forward:	9.40M cycles
Sum array backward:	9.63M cycles
Sum array even/odd:	9.64M cycles
Sum array random:	36.29M cycles
Sum array fwd ind:	11.89M cycles
Sum array unrolled:	4.78M cycles

Q: Why??

A: THE MEMORY HIERARCHY

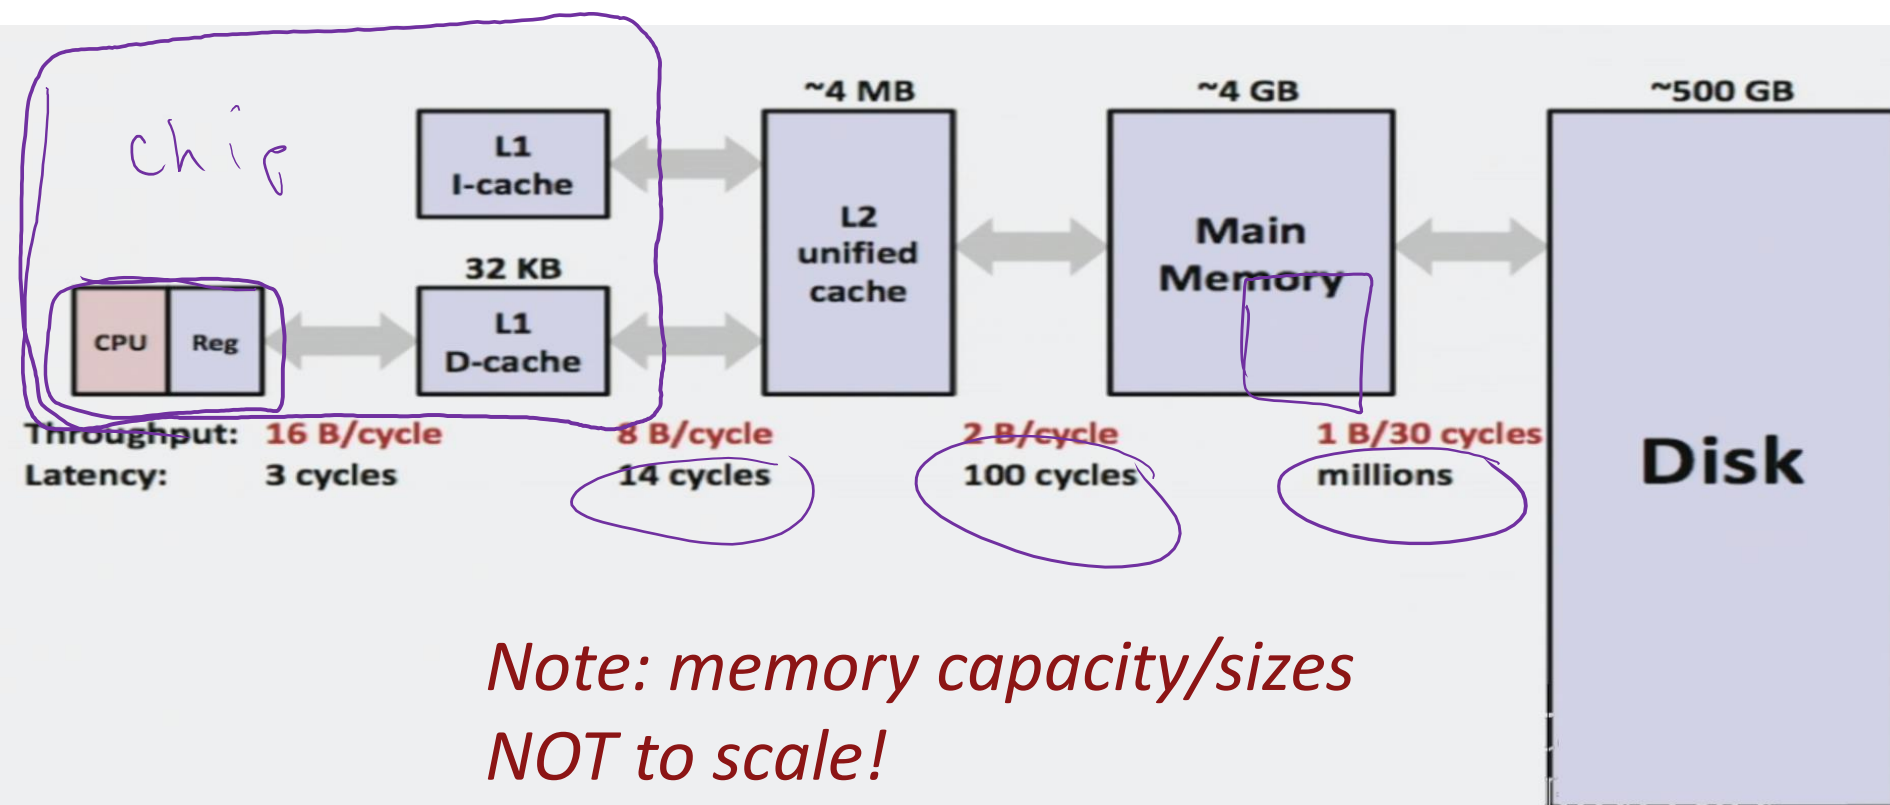
Why?? The memory hierarchy



Cache Performance Goal


- Goal:
 - › We want to have the entirety of main memory available to the ALU to perform operations at the speed of register access
- Reality says: LOL / “We all want things”
- It’s partly the cost that is a barrier (smaller/faster memory is more expensive), but there are also physical limitations

The memory hierarchy




Mapping the Memory Hierarchy


Growing gap between processor speed and memory access speed



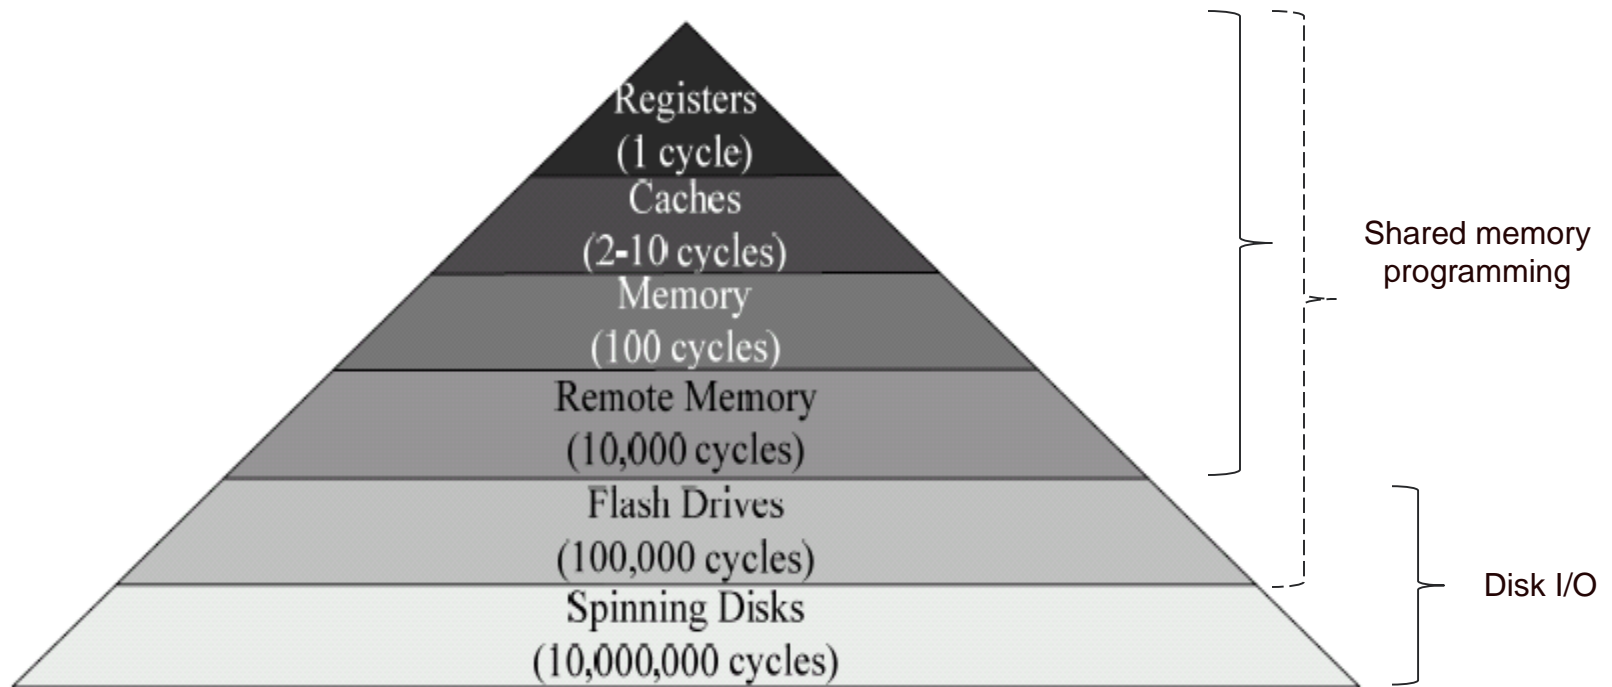
Memory Map of a Compaq Alphaserwer



Memory Map of an IBM SP-3



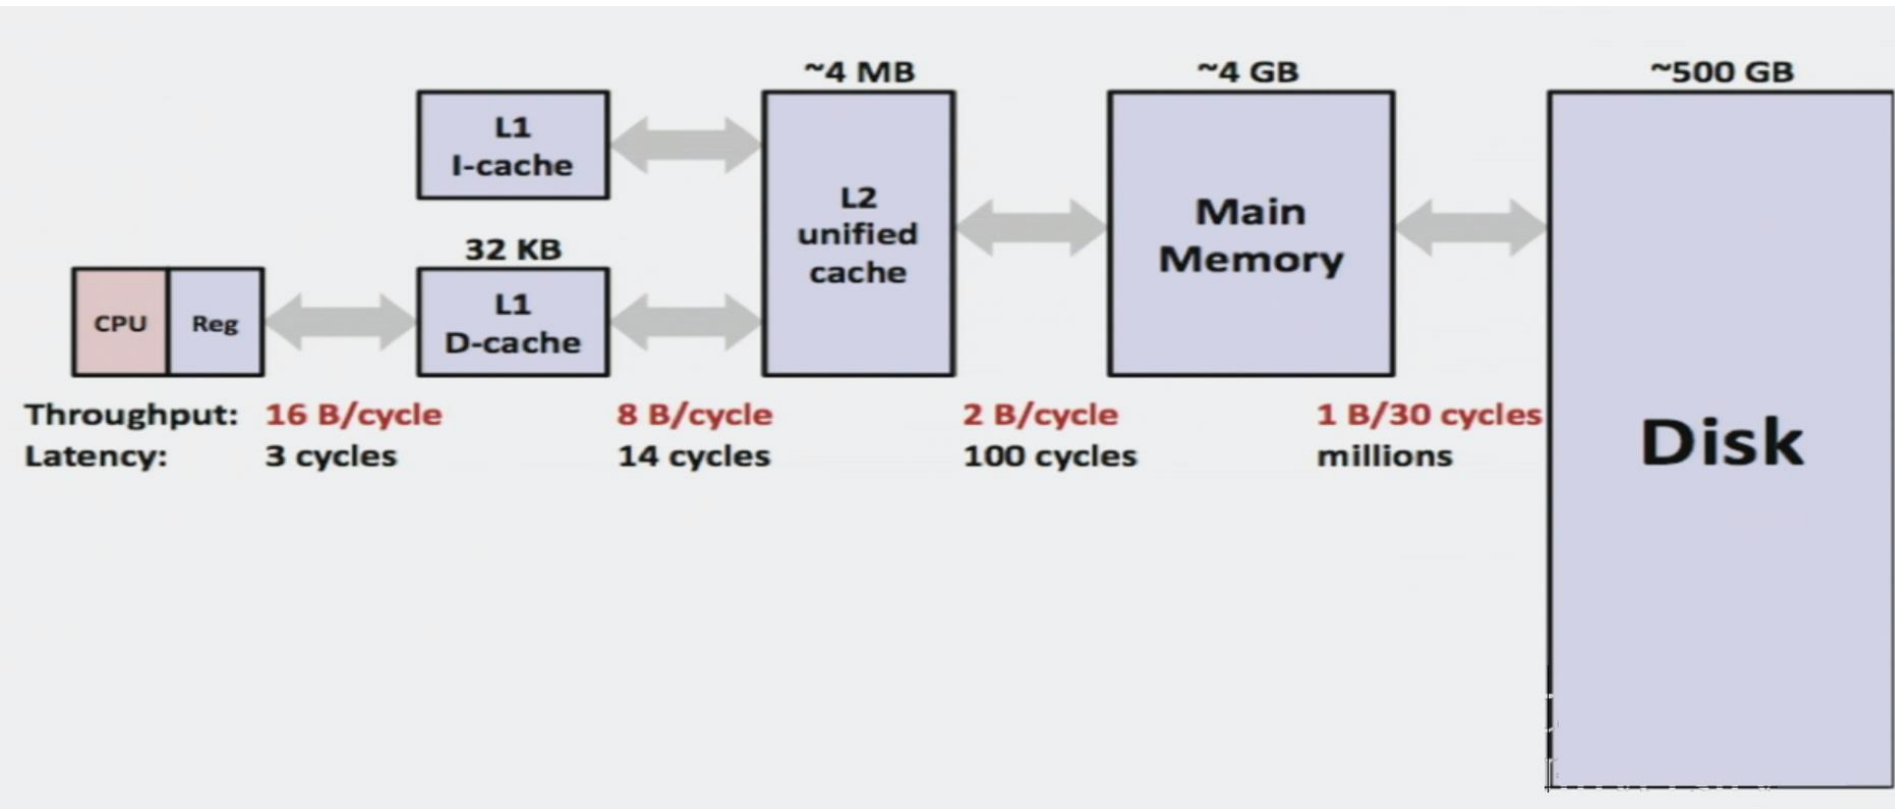
The Memory Hierarchy of Gordon (a large “supercomputer”)



Caching basics

Principles of caching

The memory hierarchy



All of caching relies on LOCALITY

- **Temporal locality:**

- › Uses of the same piece of data tend to be near each other in TIME
- › Things I have recently used, I am more likely to use again
- › *Ugliest shirts that I should probably just give away are at the very bottom of my shirt drawer, because I never wear them*

- **Spatial locality:**

- › Uses of pieces of data tend to be near each other in SPACE
- › Even if I have never used something, if it is near a used item, it is more likely to be used soon
- › *Coat closet gets ignored during summer, but once autumn hits and I use one item from there (scarf) it is likely I will soon use other items from there (various jackets)*

Cache outcomes

- **Cache hit:**
 - › What I wanted is in cache—lucky me!
 - › Hit rate: % of accesses that are cache hits

- **Cache miss:**
 - › What I wanted is not in cache—sad times!
 - › Go to main memory (or lower level of cache) to get the item, will take much longer

Example: cache performance

- Pretty realistic scenario:
 - › 97% cache hit rate
 - › Cache hit: 1 cycle to access
 - › Cache miss: 100 cycles to access
- **What percent of your total memory access time is spent on the 3% of memory accesses that are cache misses?**
 - A. $\leq 3\%$
 - B. 30%
 - C. 50%
 - D. $> 50\%$
- › Bonus discussion: How much does this change if your cache hit rate goes up slightly to 99%?