

# Announcements

## ■ Assignment 1 Out

- Your first assignment is out and comes due this Wednesday, just before midnight.
- You may not use any late days on this first assignment, since we want to grade them as soon as possible and get you feedback so you don't make the same mistakes on Assignment 2.

# Announcements

## ■ Next week

- Come next week, we'll start up on some awesome, new material.
  - We'll discuss exceptional control flow and how it can be used to support asynchronous interrupts (so that, say, your mouse can signal the mouse driver when and only when it moves), faults (e.g. when, say, a virtual address in your stack frame isn't yet mapped to physical memory, or when you dereference **NULL** because you still think that's okay for some reason ☺), traps (which are intentional calls in exceptional circumstances because the normal flow of code needs to halt while some system service runs to open a socket, create a new executable, or exit the program), and aborts (three guesses what those do).
  - The exceptional control flow material will also grant us the opportunity to write programs that spawn off and interact with other executables. It'll be our first foray into concurrency and parallelism, how it works, and what the OS needs to do to support it. It's also fascinating to know that it's all possible in the first place.

# Today's Lecture

## ■ Agenda

- I need to work through the filesystem API examples I posted this past Wednesday
- I want to take a second pass at the design and implementation of the Unix v6 filesystem, so that we're clear how blocks are the building blocks (har) of inodes, and that each and every file maps to precisely one inode. I want you to understand what an 18-byte file, a 2048-byte file, and 34471936-byte file looks like in memory.
- I also want to be clear how blocks work differently for us depending on whether their contents contribute to a regular file or to a directory. I'll finally want to draw the state of a small Unix v6 filesystem that contains a very small directory structure, where all leaf directories contain 3 regular files.
- Work through the **copy** and **t** examples posted two days ago on Wednesday. These two examples introduce the file descriptor and how it can be used to perform low-level I/O on files (and, eventually, over network connections).
- We'll also spend 10 or 15 minutes talking about how Domain Name System (otherwise known as DNS) lookup works. I do so for two reasons:
  - DNS is a simple, fault-tolerant, distributed system that's easy to explain.
  - There are a good number of features common to both a Unix v6 file system and a DNS. That there are so many architectural similarities is no accident. Both rely on a hierarchy of human-oriented names (/afs/ir.stanford.edu/class/cs110/ and graph.facebook.com, for instance) that map to more computer-palatable names like device 1, inumber 49990, blocks 388821, 338291 for the path, 31.13.75.17 for domain name.

- Realistically, I'm probably not going to get to the DNS material before Wednesday, but I can dream we'll get to it today.
- The fact the Unix v6 filesystem's design and DNS's architecture rely on naming and name resolution speaks to the widely-held belief that **naming in computer systems** is a **fundamental principle in system design**.