

## Scheduler Activations

---

Two comparisons:

- 1) threads vs. process
- 2) kernel-level vs. user-level threads

1) kernel-level threads vs. process

- \* what is a process?
  - 1) set of instructions
  - 2) some state (memory, file descriptors, signal handlers, etc.)
- \* In Linux: `clone(SIGCHLD, 0);`
- \* what is a thread?
  - \* really just a process that shares state with another process
  - \* shares: memory, file descriptors, signal handlers
- \* In Linux: `clone(CLONE_VM | CLONE_FS | CLONE_FILES | CLONE_SIGHAND, 0);`

2) kernel-level vs. user-level threads

- \* pretty simple:
  - \* kernel threads are implemented by OS
    - use system calls to create, manage, destroy, etc.
  - \* user-level threads are implemented by some user code
    - just function calls to create, manage, destroy, etc.
    - paper mentions that the compiler can help here with register allocation. how is that? A: we don't have to save every register, just the ones actually being used, if the compiler tells us which they are.
- \* key difference:
  - OS can make fully informed decisions when scheduling kernel-level threads since it has a global view of processor resources.
  - The user-level threading library obviously cannot. At least not without scheduler activations!

3) What can go wrong when user-level threads on top of kernel-level threads?

1. Recall the lock problem.
  - a. setup:
    - processors: p1, p2
    - user-level threads: u1, u2
    - kernel-level threads: t1, t2, t3 (some other proc)
    - u1 and u2 share lock at some point
  - b. user-level schedules u1 on t1, u2 on t2
  - c. kernel schedules t1, t3
  - d. u1 takes lock
  - e. kernel deschedules t1, schedule t2
  - f. u2 tries to take lock; can't. waits and waits.
  - problem: kernel has no clue about critical sections.

1:1 systems vs. N:1, vs. M:N systems

- \* 1:1 uses one kernel thread per thread
- \* N:1 puts N user threads on one thread, and uses just one thread
  - example: event-like things: Node.JS
- \* M:N puts M user threads on N kernel threads
  - common user-level threads, go-routines

---

## Scheduler Activations

Effective Kernel Support for the User-Level Management of Parallelism

Thomas Anderson, et. al

SOSP 1991

Grand Idea

- Kernel processes are:
  - great because the kernel knows about them
    - so can do a relatively good job of scheduling them.
  - not great because they tend to be slow (i.e, context switching)

- Kernel-level threads are:
  - good for the same reasons as kernel processes
  - still pretty slow (10x slower than user-level threads, according to paper)
- User-level threads are:
  - great because they tend to be fast (i.e., context switch via function call)
  - not so great because the kernel has no idea they exist
    - many user-level threads on one process
    - kernel may reschedule an important user-level thread, etc.
- So, give user-level threads an OS API to make them better
  - somehow have the kernel know about what's going on with threads
  - somehow have the user-level thread tell the kernel what's going on
  - the mechanism to do both is called a "scheduler activation"

#### Effective Kernel Support for User-Level Management of Parallelism (3)

- kernel provides user-level library with its own virtual multiprocessor
- here are the key aspects to this abstraction:
  - kernel may change number of processors in that multiprocessor
  - user-level library controls which threads to run on those processors
  - kernel notifies user-level library when:
    - an event it occurs: the user-level thread does the actual handling
      - 1) it changes number of processors
      - 2) user-level thread blocks or wakes up in the kernel
  - U-L library notifies kernel when it wants more/needs fewer processors
    - U-L library only notifies kernel of things that affect processor alloc.

#### Kernel->User Communication (3.1)

- mechanism for this U-L to K-L communication is called "scheduler activation"
  - name chosen because K-L events activate U-L library's scheduler
  - scheduler activations serve three roles: they
    - 1) are the execution context for running user-level threads
      - just like a kernel-level thread: many U-L threads per activation
      - it seems.
    - 2) notify the user-level thread system of a kernel event
    - 3) provides space in kernel for saving U-L thread context when blocking
- scheduler activations look pretty similar to traditional kernel threads
  - contains two execution stacks: one for kernel, one for app
  - U-L thread scheduler runs on the activation's user-level stack
  - each U-L thread is allocated its own stack when it starts
- when a program starts:
  - kernel creates a scheduler activation, assigns it to processor
  - then, upcalls into a fixed entry point
  - U-L library initializes itself and runs the main app. thread
- main thread may request more processors
  - kernel creates an additional scheduler activation for each processor
  - upcalls into the U-L to tell it that new processor is ready
- when kernel needs to notify U-L of events
  - create new scheduler activation, assigns it to processor
  - jumps (upcalls) into some entry point
  - then, app can do whatever it wants, just like if it was in a k-thread
- crucial distinction is:
  - kernel never resumes stopped (i.e., because of blocking) U-L threads
  - instead, new scheduler activation is created
    - notifies U-L of stopped U-L thread
  - U-L decides what to do by:
    - saving state and "removing" old thread for old activation
    - tells kernel old activation can be reused
    - decides which thread to run on the processor
- invariant: # RUNNING scheduler activations == # virtual processors
  - when new processor is added, new scheduler activation
  - because when U-L blocks, launch new scheduler activation
- the following are the scheduler activation upcall points:
  - 1) new processor added
  - 2) processor has been preempted
  - 3) scheduler activation has blocked
  - 3) scheduler activation has unblocked

- usually occur in combinations, so only one scheduler activation is created
- Example: a user-level thread blocks in the kernel
  - 1) kernel uses a fresh scheduler activation to notify U-L of event
    - so this means you have a runnable scheduler activation (this one!)
    - U-L can run other U-L threads on this activation
  - 2) when U-L thread unblocks, kernel uses fresh activation to notify U-L
    - remember: invariant of 1 scheduler activation per "processor"
    - if U-L has no "processors" (how?) kernel allocates new processor
    - notifies U-L of new processor AND of resumable blocked thread in one
    - if U-L has "processors", may have to preempt a "processor" to do upcall
    - first notifies of resumable blocked thread
    - then of preempted thread (which was on the preempted "processor")
    - U-L can decide which of the two (preempted or resumable) to execute
    - to resume a U-L thread, need that thread's state
      - most of the state is already in U-L: stack, control block
      - registers are saved by the kernel on blocking call
      - kernel passes registers to U-L when notifying of I/O completion
- Example: kernel needs to take a processor away from U-L A and give to B
  - 1) interrupts processor in A, stops existing activation
    - kernel doesn't need permission to steal processor
  - 2) "moves" processor to B, does upcall on this process in B w/new activation
  - 3) notifies A on existing some existing processor on a
    - starts new activation by preempting whatever's on that processor
    - notifies about two preemptions:
      - 1) moving processor away
      - 2) the old activation on this processor
  - 4) U-L A decides what to do with the two preempted U-L threads
    - what happens when last processor is moved away?
      - delay notification until kernel eventually re-allocates it a processor
- U-L library might want to maintain priorities of U-L threads
  - so, when a higher priority thread gets preempted, U-L can
    - somehow ask kernel to preempt the processor with the lower priority
    - can do this because U-L knows where threads are
- scheduler activations can be used to implement things other than U-L threads
  - kernel doesn't need to know about the data structures used to represent parallelism at the user-level
- scheduler activations work even when there are no U-L threads
  - IE, when only the U-L thread manager is running
  - simply creates a new scheduler activation
  - reentrant U-L thread managers can then do what they need in there
- if a U-L thread needs to do work in the kernel after some I/O unblocks
  - kernel does the right thing and does the work, THEN upcalls U-L

#### User->Kernel Communication (3.2)

- the U-L thread system need not tell the kernel about every operation
  - key observation
- U-L system notifies the kernel when:
  - 1) it has more runnable threads than "processors"
    - "add more processors (#)"
  - 2) it has more "processors" than runnable threads
    - "this processor is idle"
- Two cases to look at:
  - a) U-L system has more runnable threads than "processors"
    - if kernel doesn't assign new processors, then they MUST be busy
  - b) U-L system has more "processors" than runnable threads
    - if kernel doesn't remove unused processors, then system must be idle
    - so, U-L can keep adding parallelism without notifying kernel
- these notifications are only hints
  - requesting a "processor" doesn't guarantee you a processor NOW
- requests are serialized
  - U-L requests processor
  - ... time passes ...
  - U-L finally gets processor, but doesn't need it
  - U-L must say it doesn't need it
    - Uh...so you have to trust them?

- Yeah, but not unique to this system.
- Yeah, kernel doesn't actually know about what's running.
- Can incentivize honest guys.

#### Dealing with preemption during critical sections (3.3)

- don't want bad stuff to happen when a U-L thread gets preempted
  - like scheduling a U-L thread that's just going to sit there and wait
  - or worse, dead-lock:
    - preempted thread holds lock on user-level thread ready list
    - newly scheduled thread marks preempted thread a RUNNABLE and tries to put preempted thread into list (this sounds like bad programming)
- two options:
  - 1) let the kernel know that you're entering a critical section
    - slow.
  - 2) have a way to recover when this happens
    - this is the approach they take
- U-L system will receive an upcall when a thread is preempted
  - U-L can check if the preempted thread was in a critical section
    - if so, U-L system continues that thread until it's out of the critical section
  - once critical section is done, U-L can safely place preempted thread on ready list
  - Q: fun question: what if that U-L gets preempted while waiting?
    - then just do the same thing in the activation: stack them up

#### Implementation

- + 1200 lines to kernel (Topaz), + ~250 to user-level library (FastThreads)
  - most code was concerned with implementing processor allocation policy
- Processor allocation policy tries to be fair:
  - tries to never waste processors
  - distributes free processors evenly across spaces that need some
- Their implementation also has regular kernel threads
  - internal kernel implementation uses scheduler activations
    - so they actually do internal "downcalls", but, you know, just f-calls

#### Performance Enhancements (4.3)

- So, you want this to fast.
- most performance consideration pertain to dealing with critical sections
- somehow, U-L systems needs to know that a U-L thread is in a critical sect.
  - one way is to set a flag when it enters and clear when it leaves
    - apparently, this is not fast enough
      - adds ~10us, which is ~20% overhead (from 5.1)
    - ideally, don't want to do any work unless preemption occurs
- instead, use the compiler (lol)
  - mark each critical section in code
    - so know you know the start and end PC of the critical section
    - what about function calls in there?
      - they use a flag in this case
      - hopefully is rare. try to inline? make them small?
  - make copy each critical section
    - add a few lines to end to yield processor back to resumer
      - IE, the U-L scheduler thing
  - when U-L system gets preemption upcall
    - check PC to see if in critical section or if the flag was set
    - if in critical section, jump to the copied code
    - copied code will jump back to the resumer
      - how? probably just does a 'ret', so need to do 'call' to copy
- scheduler activations are kinda expensive to create
  - so cache their structures for reuse (like a slab allocator, sounds like)
- discarded scheduler activations can be collected and returned in bulk
  - so batching