

NFS

===

We are going from threads to distributed systems and back again
Both topics are squarely in scope at OS conferences
Similar kinds of reasoning apply to parallelism & distributed systems
E.g., Already saw happens before applied to race detection
Different "gotchas" in each case
Parallel systems: weak memory consistency can be very unintuitive
Distributed systems: node & network failures complicate reasoning
If you like distributed systems, consider CS244b next Fall...

How did Unix systems access network storage before NFS? ND - network disk
Looks like a disk device on client (e.g., /dev/nd0)
IO requests go over net to server, which has file storing disk image
Note concept still exists in the form of NBD under linux

What were goals of NFS? (p. 119)

1. Machine and OS independence
Required adding a new system call, getdirentries (p. 124). Why?
Previously, readdir used read(2) to look at raw directory contents
getdirentries abstracts directory formats across machines
2. Crash recovery
When is a local file system allowed to lose data in a crash?
Data written recently (<30 sec ago), and no call to fsync
Okay because processes writing files are killed in crash, too
When should NFS be allowed to lose data?
Client crash? Same as local file system
Server crash? Never, because processes on client not killed
3. Transparent access
"Programs should not be able to tell whether a file is remote or local."
What would be alternative?
Hack libc so open("/nfs/server/file") doesn't do real open syscall
Would break applications--e.g., inheriting file descriptors
4. UNIX semantics maintained on client
Some examples of tricky UNIX semantics?
Most permission checks happen when file initially opened
Can delete an open file and still use it
5. Reasonable performance
No worse than ND, or 80% as good as local disk

Which of above goals met by ND? All but #1

Why isn't ND good enough?

NFS also wants *sharing* of file system resources
Can't share a read-write file system stored on ND
Ordinary file systems don't expect disk to change out from under them
Might cross-allocate blocks/inodes, create duplicate file names, etc.

What does NFS look like to administrator?

Unix namespace comprised of *mount points*

A root file system mounted on "/" (the root directory) at boot time
Other file system mounted on other directories
Created with mount command: mount /dev/sda3 /usr

Server admin can *export* some file systems over NFS in /etc/exports

Client admin can mount remote FS "server:/dir" instead of device

E.g.: mount my-server:/disk/ul /home/ul

Now /home/ul on client is same FS as /disk/ul on server!

Two kinds of NFS mount: hard and soft (p. 124)--what's the difference?

Hard mount hangs forever if server crashes

With soft mount, syscalls eventually return error if no server reply

Note, more recently NFS also offers intr vs. nointer
Kernel assumed disk requests fast, so not FS syscalls interruptible
So couldn't Ctrl-C process if accidentally accessed unavailable server
Took many years to fix this by adding intr option

Until NFS, only one kind of file system in kernel. How to abstract?
Two new abstractions: VFS, and vnodes (p. 123)
Ersatz C++ abstract classes, implemented with C function pointers
Have one VFS (virtual file system) struct per mount point
 unmount(), root(), statfs(), sync()
Have one vnode for each active inode
 lookup, open, close, create, rdwr, inactive, ...
Also contains pointer to its VFS struct

Rewrote namei (p.124) [routine mapping (dir vnode, path) -> vnode] Why?
Only the client knows which directories are mount points
Hence, cannot have server translate more than one directory at a time
So namei must walk file system one vnode loopup() at a time

VFS+vnodes routines for NFS make RPCs to server--see pp. 120-121

What happens if I say "cat dir/file" on an NFS file system?
lookup (fh1, "dir") -> fh2
lookup (fh2, "file") -> fh3
read (fh3, [offset] 0, [len] 8192)
read (fh3, [offset] 8192, [len] 8192)
...

Now what if I "cd dir; cat file" while contents cached?
(fh2, "file") -> fh3 might already be in name cache
stat(fh3) -> attr
Compare mtime/ctime in attr to cached version. Only read if no match.

p. 120: "NFS uses a stateless protocol"--what does this mean?
Obviously a file system cannot be stateless
But some protocols (e.g., CIFS, 9p) keep client state on server
E.g., each file open, close sent to server
Server keeps equivalent of file descriptor table for each client
Remote descriptors maybe invalidated by server reboot/network outage
May be hard for client to rebuild state if files renamed
Or... client crashes and server stuck keeping useless descriptors
NFS requests are self-contained; no per-client context on server

What happens if server crashes?
Client keeps retransmitting requests until server answers
Once server reboots, will eventually get client's request
Because "stateless", server can execute request with no prior context

What happens if client->server request dropped by network?

Client keeps retransmitting (over UDP), so no issue

What happens if server->client reply dropped by network?

Client keeps retransmitting (over UDP), sends second copy of request

What does server do upon receiving second copy of a request?

Each RPC starts with 32-bit XID

Server keeps map of XID->reply message in a *replay cache*

Reply to duplicate requests from replay cache, rather than reexecuting

But the cache is state... protocol not stateless after all!

What happens if server->client dropped and server reboots?

Reboot wipes server replay cache, so server reexecutes request

Saving grace: protocol makes operations as *idempotent* as possible

Means executing twice has same effect as executing once

Example: "x = 1" is idempotent; "++x" is not

Of calls on pp. 120-121, which are *not* idempotent?

Not idempotent: remove, rename, link, mkdir, rmdir

create is idempotent (makes exclusive create unreliable)

Spec suggests symlink idempotent (maybe no error if content matches)

What if I say "echo message > file" on NFS file system?

```
create (fh1, "file") -> fh2
write (fh2, 0 [offset], "message") -> 2
```

What if server crashes during write?

Okay, client will keep retransmitting until it gets reply

What if server crashes after replying to write?

Client will stop retransmitting (since it gets write reply)
Means server cannot reply to write until data safely on disk

Why did authors need to work on IP fragmentation code?

Maximum Ethernet packet size is 1500 bytes [no jumbo frames in 1984]

Each request/response constitutes one UDP packet

What if you break large sequential writes into ~1400-byte requests?

Smaller than FS block size.

May have to read surrounding block

Will likely lock buffer when first write request comes in

Because writes synchronous, have to send to disk

Almost certainly pay full disk rotation before servicing next write

Solution: up to 9000-byte UDP packets (8192 data plus some header)

Must be broken into multiple IP fragments to send over Ethernet

Even back-to-back 8KiB synchronous writes likely to be slow

For good performance, want more concurrent write requests at server

What do they do? Added block I/O daemon on client (p. 125)

Not really good support for "asynchronous RPC" in kernel

Instead fork off 4-16 block I/O daemons (biode)

On write system call:

If biode is available, hand it request return from syscall immediately

biode will keep retransmitting until it gets a reply

On fsync (and later close)

Block waiting for all write requests to go through

Server hack: nfsd daemon makes system call that never returns--why?

Want to handle multiple incoming requests concurrently

Again, async IO not so easy, and want to integrate with scheduler

Easiest way to integrate with scheduler to be a process

But NFS implemented in kernel, so just do it all in a big syscall

What happens when you run "mount server:/dir /mnt"?

User-level mount program executes MOUNTPROC_MNT("/dir") to server

User-level mountd program on server returns 32-byte NFS file handle

mount program makes mount(2) syscall with server IP, fhandle, "/mnt"

NFS mount handler allocates VFS, hangs it on /mnt's vnode (in root fs)

What's in an NFS file handle?

Whatever the server wants--it's opaque to the client

Must uniquely identify file, so:

- File system ID (new field in superblock)
- File system major/minor device number (usually)
- Inode number
- Generation number (new field in inode)

What's the point of a generation number?

Generation number changes each time inode recycled

Say server deletes a file while client has it open

If client uses old handle, generation number wrong, gets ESTALE

Not great, but better than reading/writing to unrelated file!

Also makes file handles hard to guess (access control at mount time)

What is the security model? Who enforces what permissions?

Assume numeric user/group IDs are same on client and server

On server, mountd restricts permitted clients (/etc/exports)

If client sends valid file handle, NFS server assumes it is authorized

Client tags each client RPC with local user/group IDs

Server enforces access control based on claimed credentials

But maps root to -2

Does this pose problems for UNIX semantics (goal #4)?

UNIX does permission only at file open. How does NFS handle:

- fd = creat("lockfile", 0444); write(fd, ...); close(fd);
Server allows write if owner of file matches, even if 0444 perms
- fd = open(...); setuid(pw->pw_uid) /* drop privs */; read(fd, ...);"
Client sends credentials from when file originally opened (p. 127)

How is evaluation?

Close-to-open consistency