

## Read-Log-Update

=====

What is Read-Copy-Update (RCU)?

An API for synchronizing access to read-mostly data structures

Basic API for reader:

```
rcu_read_lock()    - start read-side critical section (can nest)
rcu_dereference()  - get RCU-protected global pointer
rcu_read_unlock()  - end read-side critical section
```

Basic API for writer:

```
rcu_assign_pointer() - set RCU-protected pointer to new value
synchronize_rcu()    - wait for all readers to be done with old ptr
call_rcu()            - register async. callback for when readers are done
```

Example:

```
Global: global_obj *gp; spinlock gp_lock;
Reader: rcu_read_lock(); use(rcu_dereference(gp)); rcu_read_unlock();
Writer: lock(&gp_lock);
        global_obj *old = gp;
        rcu_assign_pointer(gp, alloc_updated(old));
        unlock(&gp_lock);
        synchronize_rcu();
        obj_free(old);
```

What guarantees does RCU make?

See <http://www.efficios.com/pub/rcu/urcu-main.pdf> (Desnoyers) fig. 2

A read-side critical section can overlap at most two of following in writer:

```
+-----+-----+-----+
| assign_pointer | grace period | free object |
+-----+-----+-----+
```

How would you implement RCU? Multiple strategies:

1. Disable preemption in readers, run on all CPUs in synchronize  
rcu\_read\_lock() disables preemption  
rcu\_dereference() is memory\_order\_consume load  
rcu\_synchronize() { for (i = 0; i < ncpus; ++i) run\_on\_cpu(i); }
2. Periodically call rcu\_quiescent() in readers not in critical section  
\_Atomic uint64\_t rcu\_ctr; mutex\_t rcu\_lock;  
thread\_local \_Atomic uint64\_t local\_ctr;  
void rcu\_synchronize() {  
 lock(&rcu\_lock);  
 local\_ctr = ++rcu\_ctr;  
 for\_each\_thread { while (thread's local\_ctr < local\_ctr) yield(); }  
 unlock(&rcu\_lock);  
}  
void rcu\_quiescent() {  
 local\_ctr = global\_ctr;  
}
3. Use fences or synchronized atomics inside rcu\_read\_lock/rcu\_read\_unlock  
rcu\_read\_lock() looks like rcu\_quiescent() which no longer exists  
Last rcu\_read\_unlock() sets local\_ctr = MAXINT

Note you don't have to copy the whole data structure with RCU

Can you remove item from linked list without copying list?

Yes. See RLU Figure 1 and Desnoyers Figure 3

What happens if reader traverses linked list forwards & backwards?

Writer can't atomically swap two pointers simultaneously

Risk seeing invariant violations like item->next->prev != item -- oops

So works for some structures (singly-linked list) & not others (doubly)

What about concurrent updates? Also quite dangerous

The upshot: Need to be very clever for each data structure

Goal of read-LOG-update (RLU) is to fix all of this

Instead of swapping one pointer atomically, swap many items atomically

Maybe even atomically make updates by multiple concurrent writers!

If every problem in CS can be fixed with a level of indirection

... what's the magic level of indirection in RLU?

We will actually update a bunch of data structures non-atomically

But, some GLOBAL CLOCK makes readers ignore all of these updates  
Then atomically update the clock and new readers will see new values  
In-progress readers will continue to use old clock value!

What is RLU API?

`rlu_reader_lock()` / `rlu_dereference(obj)` - like RCU  
`rlu_reader_unlock()` - Like RCU, but might also commit writes you made  
`rlu_try_lock(ptr)` - May return locked private object copy you can modify  
`rlu_cmp_objs(ptr, ptr)` - because unlike RCU "`obj1 == obj2`" may not work  
`rlu_assign_pointer(lptra, ptr)` - deals with multiple pointer values properly  
`rlu_abort()` - abort current transaction and undo any pending writes  
`rlu_malloc()/rlu_free()`? (Reserve space for `log_ptr`)

What data structures do we need (Sec. 3.4)?

- \* Per thread:
  - 2 Write logs: one active, one left over from previous epoch
  - `run_counter`: odd when using RLU, even when outside critical section
  - `local_clock`: Time at which your view of data should be frozen
  - `write_clock`: Infinity until you want to commit
  - `is_writer`: Flag is true if you have made modifications
- \* Global:
  - `global_clock`: Used to initialize local clock each transaction
  - List of threads providing access to per-thread info for all threads
- \* Per object in heap
  - `log_ptr`: pointer to copy of object in a log (or NULL if none)
- \* Per object in log
  - Thread identifier of thread that owns log
  - A pointer to the original object of which this is a copy
  - The size of the object (so you know where next log entry is)
  - Special reserved value (where `log_ptr` would be in heap object)

How does RLU work with serial writers (assume external big write lock)?

For reader:

`rlu_reader_lock()`

- bump `run_counter` to something odd
- `local_clock` = `global_clock` - this determines version of data you see
- `is_writer` = false

`rlu_reader_unlock()` - bump `run_counter` to something even  
`rlu_dereference(ptr)`

- unlocked heap object? return as-is
- log object? return as-is (you had the pointer, must be okay)

Otherwise, we must have a locked heap object

- Did we lock it? Then return corresponding log object
- Is `local_clock`  $\geq$  log owner's `write_clock`? Then return log object
- Otherwise, locked after `rlu_reader_lock()`, so return heap object

`rlu_assign_pointer(lptra, ptr)` - always translate `ptr` to heap if log

For writer:

`rlu_try_lock(ptr)`

- Initialize log header for object copy in our log
- Set heap object `log_ptr` to point to our header
- Copy heap object into log

`rlu_synchronize()`

- Record `run_counter` of all active threads
- For each thread T, spin while all of:
  1. T's `run_counter` is odd (so T using RLU), and
  2. T's `run_counter` is the same as we recorded before the loop, and
  3. T's `local_clock` < our `write_clock` (T ignoring our log)

`rlu_reader_unlock()`

- set `write_clock` = `global_clock` + 1 (was infinity)
- ++`global_clock` - now new readers will use our log
- `rlu_synchronize()`
- set `write_clock` = infinity
- Swap our two logs (readers might still use recent one - 3.6.2)

What happens if you add concurrent writers?

We can use `log_ptr` as a per-object lock

Just modify `rlu_try_lock(ptr)`

- Already locked by us? return `log_ptr`
- Already locked by another thread? `rlu_abort()`
- Initialize log header for object copy in our log
- Try to lock with atomic compare-and-swap on heap header
- Success? Copy object from heap to log. Failure? `rlu_abort()`

Drawbacks of concurrent writer version?

Aborts (don't happen in serial writer version)

More complicated semantics with concurrent writers, e.g.:

With serial writer RLU, following code guarantees `a == b`:

T1: `reg = *rlu_dereference(a); *rlu_try_lock(b) = reg;`

T2: `reg = *rlu_dereference(b); *rlu_try_lock(a) = reg;`

With concurrent writer version, might just swap `a`'s and `b`'s values

Why do you want RLU deferring (3.7) and what is it?

`rlu_synchronize()` is expensive. When do you really need it?

After copying object into heap (to ensure old `log_ptr` not used)

Therefore if another thread wants to write same object

Also if you want another thread to see changes--why?

Say just you wrote back log object w/o incrementing counter

Then set `log_ptr` to NULL to release the lock

Now readers whose `local_clock < your write_clock` will see new version

Can't set all `log_ptr`s simultaneously - violates atomicity

A readers may see both versions in one transaction - violates isolation

How to reduce calls to `rlu_synchronize()`?

Omit bumping `global_clock` and calling `rcu_synchronize` in `rcu_read_unlock`

Instead send "sync request" to other thread after write-write conflict

Anticipate fewer sync requests than `is_writer` transactions

Bonus: reduces contention on global clock (to write-write conflicts)

Bonus 2: fewer conflict cache misses in readers

Yes, but because they are seeing stale data

Drawbacks?

Usage scenario must tolerate added update latency

What if sync request takes too long? Maybe write back other thread's log?

Only a big win with one benchmark at (Citrus to 80 cores)

More modest win with only 16 cores

What evaluation questions should be asking?

1. How is performance compared to RCU today?

On average systems today

When scaled to many cores

2. How is correctness?

Is the approach correct (these things are tricky)?

Is it easy to implement correctly, or bug prone?

3. Does it simplify/allow more data structures than RCU?

What RCU do they compare to--is this fair?

"State-of-the art" `urcu` library (Desnoyers) and linux kernel RCU

Include numbers for incorrect but faster Harris linked list w. mem. leak

Authors even fixed a performance bug in linux `list_entry_rcu` (4.6)

But... authors don't support all features of more mature RCUs

\* Fig. 4--What is the best linked list?

RLU unless you want memory leaks

Note how concurrent updates kill RCU scalability!

Don't show deferred RLU but claim it is the same

\* Fig. 5 (hash table)

Why does RCU dominate?

RCU readers do less constant work than Harris readers (and RLU)

Use RCU on individual buckets--means very few actual write conflicts

Why does deferred update win here but not in linked list?

Fewer write-write conflicts in bucketed hash table than in single list

\* Fig. 6 (resizable hash table)

How does a resizable hash table even work in RCU?

Super complicated because can only update one pointer at a time

So first create expanded hash table where two buckets share each list

Assumes you aren't completely rehashing, but just splitting buckets  
 Not useful, yet, but at least lookups in new table correct  
 Then go through and "unzip" the buckets so they share tails  
 Eventually tails will be NULL and you are done  
 How does a resizable hash table work in RLU?  
 Just like a regular hash table + some rcu\_dereference/rcu\_try\_lock  
 Which performs better RCU or RLU for resizable hash table?  
 Figure 6 appears to show slight win for RCU. Is this fair?  
 Resizable hash table doesn't support concurrent insert/remove  
 Benchmarks is from other paper, so designed to make RCU look good  
 Might well see win for RLU under different workload

- \* Fig. 7 (stress test)--Why is RLU twice as slow even for 1 CPU?  
 100% update, one item per bucket is best case for RCU  
 Cost of copying objects in RLU causes 2x slowdown  
 At least this likely bounds worst-case for RLU
- \* What do we learn from Citrus search tree (Sec 4.5)?  
 Complex structures possible with RCU, but complicated  
 Compare listing 3 to listing 4--RLU wins on simplicity  
 Plus RLU performs better

Do authors answer questions 2 and 3?

- 2 (correctness): passes applicable subset of kernel torture test
- 3 (functionality):  
 Already saw how much simpler Citrus code got  
 Kyoto cabinet Cache DB  
 RLU speeds up application where RCU is intractable