Run-time Environments

Lecture 11

Status

- We have covered the front-end phases
  - Lexical analysis
  - Parsing
  - Semantic analysis
- Next are the back-end phases
  - Optimization
  - Code generation

- We'll do code generation first . . .

Run-time environments

- Before discussing code generation, we need to understand what we are trying to generate

- There are a number of standard techniques for structuring executable code that are widely used

---

Outline

- Management of run-time resources

- Correspondence between
  - static (compile-time) and
 - dynamic (run-time) structures

- Storage organization

Run-time Resources

- Execution of a program is initially under the control of the operating system

- When a program is invoked:
  - The OS allocates space for the program
  - The code is loaded into part of the space
  - The OS jumps to the entry point (i.e., "main")

Memory Layout

1

Notes

- By tradition, pictures of machine organization have:
  – Low address at the top
  – High address at the bottom
  – Lines delimiting areas for different kinds of data

- These pictures are simplifications
  – E.g., not all memory need be contiguous

What is Other Space?

- Holds all data for the program
- Other Space = Data Space

- Compiler is responsible for:
  – Generating code
  – Orchestrating use of the data area

Code Generation Goals

- Two goals:
  – Correctness
  – Speed

- Most complications in code generation come from trying to be fast as well as correct

Assumptions about Execution

1. Execution is sequential; control moves from one point in a program to another in a well-defined order

2. When a procedure is called, control eventually returns to the point immediately after the call

Do these assumptions always hold?

Activations

- An invocation of procedure P is an *activation* of P

- The *lifetime* of an activation of P is
  – All the steps to execute P
  – Including all the steps in procedures P calls

Lifetimes of Variables

- The *lifetime* of a variable x is the portion of execution in which x is defined

- Note that
  – Lifetime is a dynamic (run-time) concept
 – Scope is a static concept

Activation Trees

- Assumption (2) requires that when P calls Q, then Q returns before P does

- Lifetimes of procedure activations are properly nested

- Activation lifetimes can be depicted as a tree

Example

```
Class Main {
   g() : Int { 1 };
   f():  Int { g() };
   main(): Int {{ g(); f(); }};
}
```

Main
g   f
           g

Example 2

```
Class Main {
   g() : Int { 1 };
   f(x:Int):  Int { if x = 0 then g() else f(x - 1) fi};
   main(): Int {{f(3); }};
}
```

What is the activation tree for this example?

Notes

- The activation tree depends on run-time behavior

- The activation tree may be different for every program input

- Since activations are properly nested, a stack can track currently active procedures

Example

```
Class Main {
   g() : Int { 1 };
   f():  Int { g() };
   main(): Int {{ g(); f(); }};
}
```

	Stack
Main	*Main*

Example

```
Class Main {
   g() : Int { 1 };
   f():  Int { g() };
   main(): Int {{ g(); f(); }};
}
```

	Stack
Main	*Main*
g	*g*

Example

```
Class Main {
    g() : Int { 1 };
    f():  Int { g() };
    main(): Int {{ g(); f(); }};
}
```

Main

g   f

Stack

Main

f

Example

```
Class Main {
    g() : Int { 1 };
    f():  Int { g() };
    main(): Int {{ g(); f(); }};
}
```

Main

g   f

g

Stack

Main

f

g

Revised Memory Layout

Low Address

Code

Memory

Stack

High Address

Activation Records

- The information needed to manage one procedure activation is called an *activation record (AR)* or *frame*

- If procedure F calls G, then G's activation record contains a mix of info about F and G.

What is in G's AR when F calls G?

- F is "suspended" until G completes, at which point F resumes. G's AR contains information needed to resume execution of F.

- G's AR may also contain:
  – G's return value (needed by F)
  – Actual parameters to G (supplied by F)
  – Space for G's local variables

The Contents of a Typical AR for G

- Space for G's return value
- Actual parameters
- Pointer to the previous activation record
  – The *control link*; points to AR of caller of G
- Machine status prior to calling G
  – Contents of registers & program counter
  – Local variables
- Other temporary values

4

Example 2, Revisited

Class Main {
  g() : Int { 1 };
 f(x:Int):Int {if x=0 then g() else f(x - 1)(**)fi};
  main(): Int {{f(3); (*)
}};}

AR for f:

result
argument
control link
return address

Stack After Two Calls to f

Main
f   (result)
    3

    (*)
f   (result)
    2

    (**)

Notes

- Main has no argument or local variables and its result is never used; its AR is uninteresting
- (*) and (**) are return addresses of the invocations of f
  – The return address is where execution resumes after a procedure call finishes

- This is only one of many possible AR designs
  – Would also work for C, Pascal, FORTRAN, etc.

The Main Point

The compiler must determine, at compile-time, the layout of activation records and generate code that correctly accesses locations in the activation record

Thus, the AR layout and the code generator must be designed together!

Example

The picture shows the state after the call to the 2nd invocation of f returns

Main
f   (result)
    3

    (*)
f   1
    2

    (**)

Discussion

- The advantage of placing the return value 1st in a frame is that the caller can find it at a fixed offset from its own frame

- There is nothing magic about this organization
  – Can rearrange order of frame elements
  – Can divide caller/callee responsibilities differently
  – An organization is better if it improves execution speed or simplifies code generation

Discussion (Cont.)

- Real compilers hold as much of the frame as possible in registers
 - Especially the method result and arguments

Globals

- All references to a global variable point to the same object
  - Can't store a global in an activation record

- Globals are assigned a fixed address once
  - Variables with fixed address are "statically allocated"
- Depending on the language, there may be other statically allocated values

Memory Layout with Static Data

Code	Low Address
Static Data	
Stack	
	High Address

Memory

Heap Storage

- A value that outlives the procedure that creates it cannot be kept in the AR

  method foo() { new Bar }

  The Bar value must survive deallocation of foo's AR

- Languages with dynamically allocated data use a *heap* to store dynamic data

Notes

- The code area contains object code
  - For most languages, fixed size and read only
- The static area contains data (not code) with fixed addresses (e.g., global data)
  - Fixed size, may be readable or writable
- The stack contains an AR for each currently active procedure
  - Each AR usually fixed size, contains locals
- Heap contains all other data
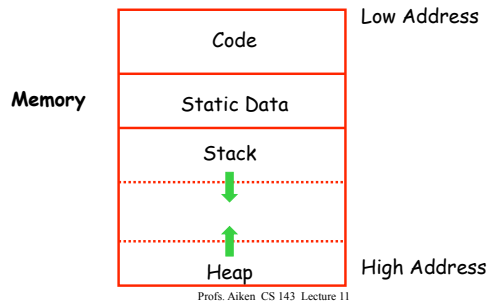 - In C, heap is managed by *malloc* and *free*

Notes (Cont.)

- Both the heap and the stack grow

- Must take care that they don't grow into each other

- Solution: start heap and stack at opposite ends of memory and let them grow towards each other

Memory Layout with Heap

Low Address

| Code |
| Static Data |
| Stack |

High Address

Heap

Memory

Data Layout

- Low-level details of machine architecture are important in laying out data for correct code and maximum performance

- Chief among these concerns is *alignment*

Alignment

- Most modern machines are (still) 32 bit
  - 8 bits in a byte
  - 4 bytes in a word
  - Machines are either byte or word addressable
- Data is *word aligned* if it begins at a word boundary
- Most machines have some alignment restrictions
  - Or performance penalties for poor alignment

Alignment (Cont.)

- Example: A string

  "Hello"

  Takes 5 characters (without a terminating \0)

- To word align next datum, add 3 "padding" characters to the string

- The padding is not part of the string, it's just unused memory


Next Topic: Stack Machines

- A simple evaluation model
- No variables or registers
- A stack of values for intermediate results
- Each instruction:
  - Takes its operands from the top of the stack
  - Removes those operands from the stack
  - Computes the required operation on them
  - Pushes the result on the stack

Example of Stack Machine Operation

- The addition operation on a stack machine

| 5 |
| 7 |
| 9 |
| … |

| 9 |
| … |

| 12 |
| 9 |
| … |

pop   add   push

7

Example of a Stack Machine Program

- Consider two instructions
  - push i - place the integer i on top of the stack
  - add   - pop two elements, add them and put
              the result back on the stack
- A program to compute 7 + 5:

  push 7
  push 5
  add

Why Use a Stack Machine?

- Each operation takes operands from the same place and puts results in the same place

- This means a uniform compilation scheme

- And therefore a simpler compiler

Why Use a Stack Machine ?

- Location of the operands is implicit
  - Always on the top of the stack
- No need to specify operands explicitly
- No need to specify the location of the result
- Instruction "add" as opposed to "add r_1, r_2"
  - ⇒ Smaller encoding of instructions
  - ⇒ More compact programs
- This is one reason why Java Bytecodes use a stack evaluation model

---

Optimizing the Stack Machine

- The add instruction does 3 memory operations
  - Two reads and one write to the stack
  - The top of the stack is frequently accessed
- Idea: keep the top of the stack in a register (called accumulator)
  - Register accesses are faster
- The "add" instruction is now

  acc ← acc + top_of_stack
  - Only one memory operation!


Stack Machine with Accumulator

Invariants
- The result of an expression is in the accumulator

- For $op(e_1,…,e_n)$ push the accumulator on the stack after computing $e_1,…,e_{n-1}$
  - After the operation pops n-1 values

- Expression evaluation preserves the stack

Stack Machine with Accumulator. Example

- Compute 7 + 5 using an accumulator

acc ← 7   acc ← 5   acc ← acc + top_of_stack
push acc   pop

A Bigger Example: 3 + (7 + 5)

Code	Acc	Stack
acc ← 3	3	‹init›
push acc	3	3, ‹init›
acc ← 7	7	3, ‹init›
push acc	7	7, 3, ‹init›
acc ← 5	5	7, 3, ‹init›
acc ← acc + top_of_stack	12	7, 3, ‹init›
pop	12	3, ‹init›
acc ← acc + top_of_stack	15	3, ‹init›
pop	15	‹init›

Notes

- It is very important evaluation of a subexpression preserves the stack
  - Stack before the evaluation of 7 + 5 is 3, ‹init›
  - Stack after the evaluation of 7 + 5 is 3, ‹init›
  - The first operand is on top of the stack